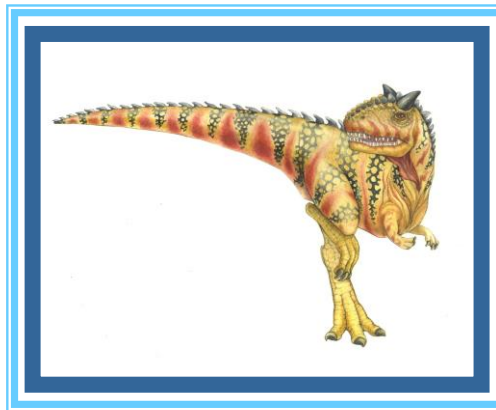


Chapter 3: Processes





Chapter 3: Processes

Process Concept

Process Scheduling

Operations on Processes

Interprocess Communication

Examples of IPC Systems

Communication in Client-Server Systems





Objectives

To introduce the notion of a process -- a program in execution, which forms the basis of all computation

To describe the various features of processes, including scheduling, creation and termination, and communication

To explore interprocess communication using shared memory and message passing

To describe communication in client-server systems





Process Concept

An operating system executes a variety of programs:

Batch system – **jobs**

Time-shared systems – **user programs** or **tasks**

Textbook uses the terms **job** and **process** almost interchangeably

Process – a program in execution; process execution must progress in sequential fashion

Multiple parts

The program code, also called **text section**


Current activity including **program counter**, processor registers

Stack containing temporary data

- ▶ Function parameters, return addresses, local variables

Data section containing global variables

Heap containing memory dynamically allocated during run time





Process Concept (Cont.)

Program is **passive** entity stored on disk (**executable file**),
process is **active**

Program becomes process when executable file loaded into
memory

Execution of program started via GUI mouse clicks, command
line entry of its name, etc

One program can be several processes

Consider multiple users executing the same program





Process in Memory

```
#include "stdafx.h"

int x;
int y = 1;

int _tmain(int argc, _TCHAR* argv[])
{
    int z;
    int *h = new int[100];

    printf("main => %p\n", _tmain);
    printf("x => %p\n", &x);
    printf("y => %p\n", &y);
    printf("z => %p\n", &z);
    printf("h is %p\n", h);

    getchar();

    return 0;
}
```





Process in Memory

```

3:
4: #include "stdafx.h"
5:
6: int x;
7: int y = 1;
8:
9:
10: int _tmain(int argc, _TCHAR* argv[])
11: {
000007F7FB402E50 48 89 54 24 10      mov     qword ptr [rsp+10h],rdx
000007F7FB402E55 89 4C 24 08         mov     dword ptr [rsp+8],ecx
000007F7FB402E59      57                push    rdi
000007F7FB402E5A 48 83 EC 50         sub     rsp,50h
000007F7FB402E5E 48 8B FC           mov     rdi,rsp
000007F7FB402E61 B9 14 00 00 00     mov     ecx,14h
000007F7FB402E66 B8 CC CC CC CC     mov     eax,0CCCCCCCCh
000007F7FB402E6B F3 AB             rep stos dword ptr [rdi]
000007F7FB402E6D 8B 4C 24 60         mov     ecx,dword ptr [rsp+60h]
12:     int z;
13:     int *h = new int[100];
000007F7FB402E71 B9 90 01 00 00     mov     ecx,190h
000007F7FB402E76 E8 DB E2 FF FF     call    operator new (7F7FB401156h)
000007F7FB402E7B 48 89 44 24 40     mov     qword ptr [rsp+40h],rax
000007F7FB402E80 48 8B 44 24 40     mov     rax,qword ptr [rsp+40h]
000007F7FB402E85 48 89 44 24 38     mov     qword ptr [h],rax
14:
15:     printf("main => %p\n", _tmain);
000007F7FB402E8A 48 8D 15 74 E1 FF FF lea     rdx,[@ILT+0(wmain) (7F7FB401005h)]
000007F7FB402E91 48 8D 0D F8 38 00 00 lea     rcx,[__xi_z+130h (7F7FB406790h)]
000007F7FB402E98 FF 15 72 86 00 00   call    qword ptr [__imp_printf (7F7FB40B510h)]
16:     printf("x => %p\n", &x);
000007F7FB402E9E 48 8D 15 AB 62 00 00 lea     rdx,[x (7F7FB409150h)]
000007F7FB402EA5 48 8D 0D F4 38 00 00 lea     rcx,[__xi_z+140h (7F7FB4067A0h)]
000007F7FB402EAC FF 15 5E 86 00 00   call    qword ptr [__imp_printf (7F7FB40B510h)]
17:     printf("y => %p\n", &y);
000007F7FB402EB2 48 8D 15 47 61 00 00 lea     rdx,[y (7F7FB409000h)]
000007F7FB402EB9 48 8D 0D F0 38 00 00 lea     rcx,[__xi_z+150h (7F7FB4067B0h)]
000007F7FB402EC0 FF 15 4A 86 00 00   call    qword ptr [__imp_printf (7F7FB40B510h)]

```





Process in Memory

```
18:    printf("z => %p\n", &z);
000007F7FB402EC6 48 8D 54 24 24    lea     rdx,[z]
000007F7FB402ECB 48 8D 0D EE 38 00 00 lea     rcx,[__xi_z+160h (7F7FB4067C0h)]
000007F7FB402ED2 FF 15 38 86 00 00    call    qword ptr [__imp_printf (7F7FB40B510h)]

19:    printf("h is %p\n", h);
000007F7FB402ED8 48 8B 54 24 38      mov     rdx,qword ptr [h]
000007F7FB402EDD 48 8D 0D EC 38 00 00 lea     rcx,[__xi_z+170h (7F7FB4067D0h)]
000007F7FB402EE4 FF 15 26 86 00 00    call    qword ptr [__imp_printf (7F7FB40B510h)]

20:
21:    getchar();
000007F7FB402EEA FF 15 28 86 00 00    call    qword ptr [__imp_getchar (7F7FB40B518h)]


22:
23:    return 0;
000007F7FB402EF0 33 C0              xor     eax,eax

24: }
000007F7FB402EF2 8B F8              mov     edi,eax
000007F7FB402EF4 48 8B CC           mov     rcx,rsp
000007F7FB402EF7 48 8D 15 22 39 00 00 lea     rdx,[__xi_z+1C0h (7F7FB406820h)]
000007F7FB402EFE E8 CD E1 FF FF     call    _RTC_CheckStackVars (7F7FB4010D0h)
000007F7FB402F03 8B C7              mov     eax,edi
000007F7FB402F05 48 83 C4 50        add     rsp,50h
000007F7FB402F09 5F                 pop     rdi
000007F7FB402F0A C3                 ret
```





Process in Memory





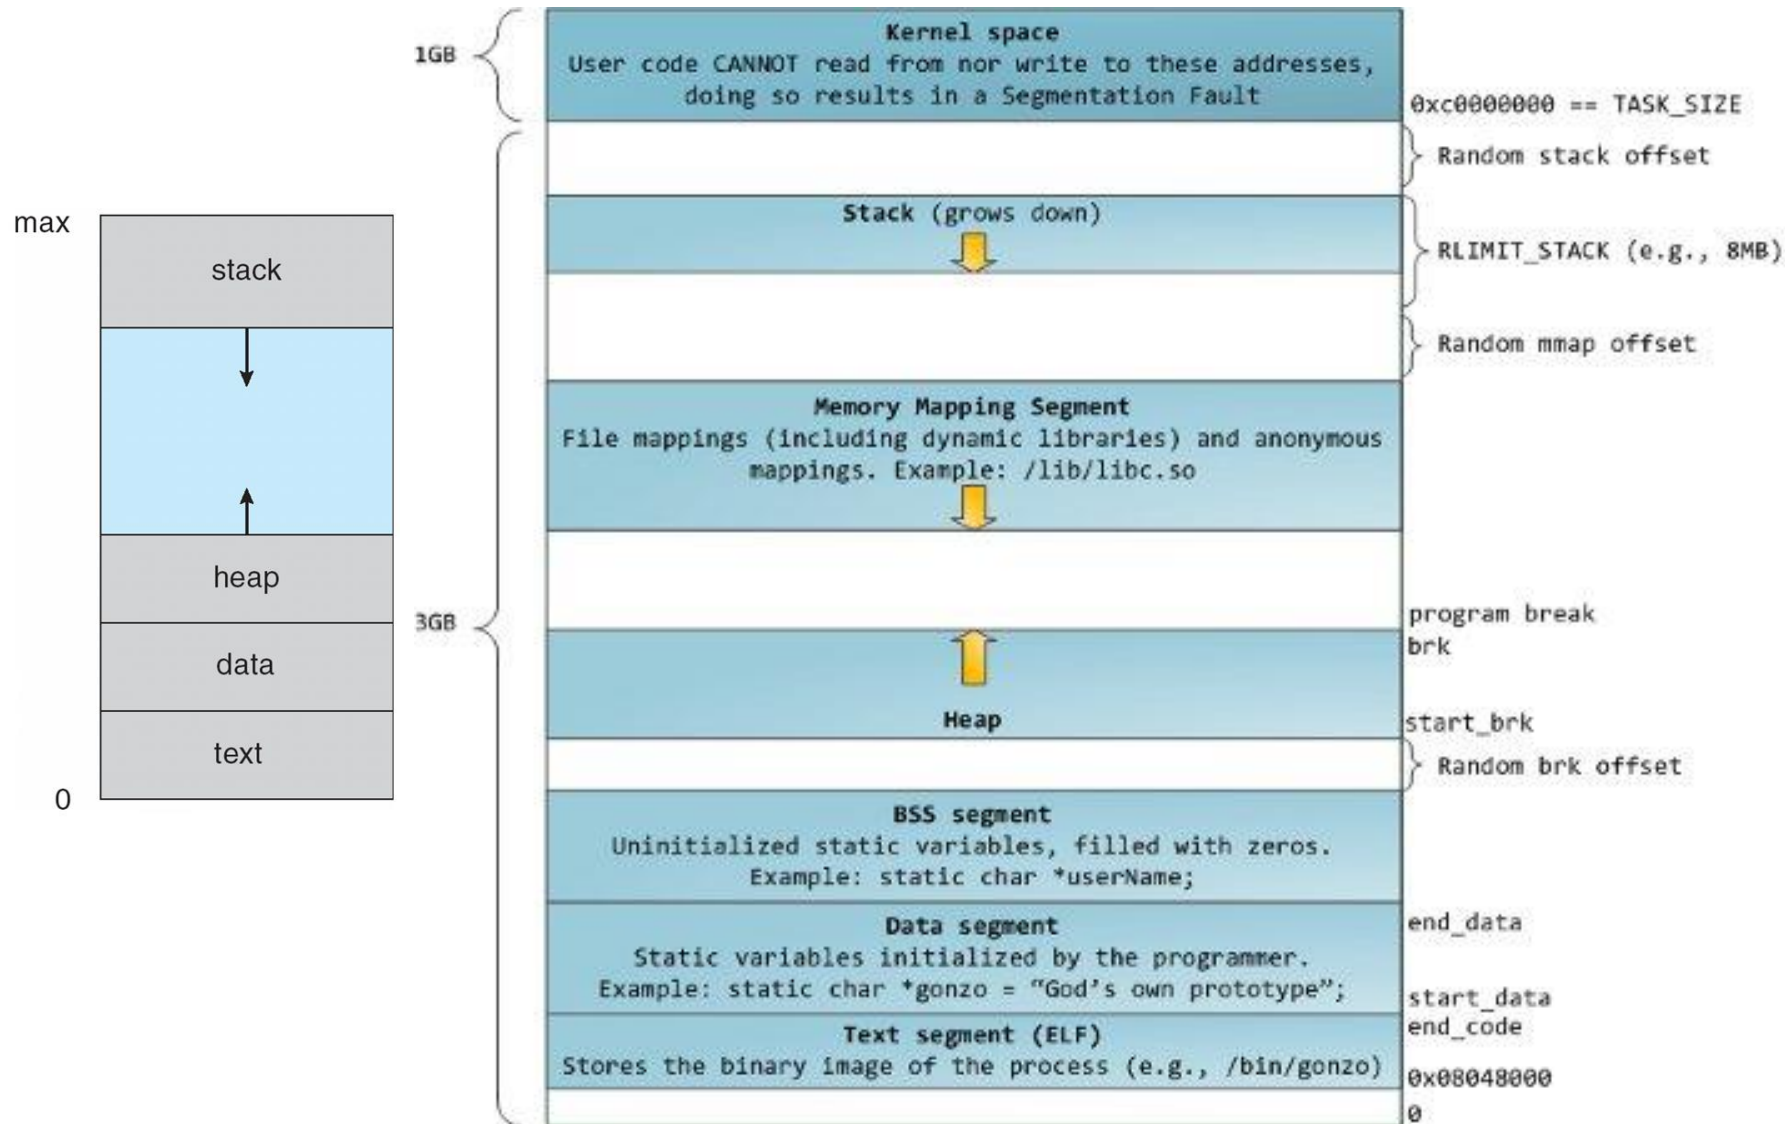
c:\users\hank\documents\visual studio 2010\Projects\test_process1\x64\...

```
main => 000007F7FB401005
x => 000007F7FB409150
y => 000007F7FB409000
z => 000000000097FAD4
h is 0000000000AD83A0
```

Address	Type	Size	Protection	Details
+ 0000000000850000	Heap (Shareable)	64 K	Read/Write	Heap ID: 2 [COMPATABILITY]
+ 0000000000860000	Mapped File	4 K	Read	C:\Windows\Globalization\zh-TW.nlx
+ 0000000000870000	Shareable	36 K	Read	
- 0000000000880000	Thread Stack	1,024 K	Read/Write/Guard	Thread ID: 1340
0000000000880000	Thread Stack	992 K	Reserved	
0000000000978000	Thread Stack	12 K	Read/Write/Guard	
000000000097B000	Thread Stack	20 K	Read/Write	
+ 0000000000980000	Shareable	16 K	Read	
+ 0000000000990000	Shareable	4 K	Read	
+ 00000000009A0000	Private Data	8 K	Read/Write	
+ 00000000009B0000	Mapped File	468 K	Read	C:\Windows\System32\locale.nls
- 0000000000AD0000	Heap (Private Data)	64 K	Read/Write	Heap ID: 3 [COMPATABILITY]
0000000000AD0000	Heap (Private Data)	36 K	Read/Write	Heap ID: 3 [COMPATABILITY]
0000000000AD9000	Heap (Private Data)	28 K	Reserved	Heap ID: 3 [COMPATABILITY]
+ 0000000000B50000	Heap (Private Data)	1,024 K	Read/Write	Heap ID: 1 [COMPATABILITY]
+ 000000005B3E0000	Image (ASLR)	1,856 K	Execute/Read	C:\Windows\System32\msvcr100d.dll
+ 000000007FFE0000	Private Data	64 K	Read	
+ 000007F7FB0F0000	Shareable	1,024 K	Read	
+ 000007F7FB1F0000	Shareable	204 K	Read	
+ 000007F7FB223000	Private Data	4 K	Read/Write	Process Environment Block
+ 000007F7FB22E000	Private Data	8 K	Read/Write	Thread Environment Block ID: 1340
- 000007F7FB400000	Image (ASLR)	56 K	Execute/Read	C:\Users\Hank\Documents\Visual Studio 2010\Projects\test.p
000007F7FB400000	Image (ASLR)	4 K	Read	Header
000007F7FB401000	Image (ASLR)	20 K	Execute/Read	.text
000007F7FB406000	Image (ASLR)	12 K	Read	.rdata
000007F7FB409000	Image (ASLR)	4 K	Read/Write	.data
000007F7FB40A000	Image (ASLR)	4 K	Read	.pdata
000007F7FB40B000	Image (ASLR)	4 K	Read/Write	.idata
000007F7FB40C000	Image (ASLR)	4 K	Read	.rsrc
000007F7FB40D000	Image (ASLR)	4 K	Read	.reloc
+ 000007FDCCDA0000	Image (ASLR)	972 K	Execute/Read	C:\Windows\System32\KernelBase.dll



Process in Memory





Process State

As a process executes, it changes *state*

new: The process is being created

running: Instructions are being executed

waiting: The process is waiting for some event to occur

ready: The process is waiting to be assigned to a processor

terminated: The process has finished execution

```
hank@Maestro:/home/hank
File Edit View Search Terminal Help
top - 23:19:48 up 3 min,  2 users,  load average: 0.56, 0.52, 0.25
Tasks: 144 total,   2 running, 142 sleeping,   0 stopped,   0 zombie
Cpu(s): 10.3%us,  7.4%sy,  0.0%ni, 81.9%id,  0.2%wa,  0.0%hi,  0.2%si,  0.0%st
Mem:   1019704k total,   878904k used,   140800k free,   34432k buffers
Swap:  2064380k total,    0k used,   2064380k free,   318172k cached

  PID USER      PR  NI  VIRT  RES  SHR  S  %CPU  %MEM     TIME+  COMMAND
 2669 hank       20   0  721m 100m  31m  S   19.3   10.0   0:15.16  firefox
 1936 root       20   0  303m  93m  14m  S   12.0    9.4   0:07.24  Xorg
 2268 hank       20   0 1836m  94m  43m  S   10.6    9.5   0:08.35  gnome-shell
 2477 hank       20   0  572m  15m 9884  S    0.3    1.5   0:00.55  gnome-terminal
 2668 root       20   0 15256 1212  900  R    0.3    0.1   0:00.21  top
    1 root       20   0 66744  24m 2080  S    0.0    2.5   0:01.20  systemd
    2 root       20   0      0     0     0  S    0.0    0.0   0:00.00  kthreadd
    3 root       20   0      0     0     0  S    0.0    0.0   0:00.01  ksoftirqd/0
    4 root       20   0      0     0     0  S    0.0    0.0   0:00.00  kworker/0:0
    5 root       20   0      0     0     0  S    0.0    0.0   0:00.00  kworker/u:0
    6 root       RT   0      0     0     0  S    0.0    0.0   0:00.00  migration/0
    7 root       RT   0      0     0     0  S    0.0    0.0   0:00.00  watchdog/0
    8 root       RT   0      0     0     0  S    0.0    0.0   0:00.00  migration/1
    9 root       20   0      0     0     0  S    0.0    0.0   0:00.00  kworker/1:0
   10 root       20   0      0     0     0  S    0.0    0.0   0:00.04  ksoftirqd/1
   11 root       RT   0      0     0     0  S    0.0    0.0   0:00.00  watchdog/1
   12 root        0  -20      0     0     0  S    0.0    0.0   0:00.00  cpuset
   13 root        0  -20      0     0     0  S    0.0    0.0   0:00.00  khelper
```







Diagram of Process State





Process State

```
hank@Maestro:/home/hank$ ps axfj | more
File Edit View Search Terminal Help

[root@Maestro hank]# ps axfj | more
PPID  PID  PGID  SID  TTY  TPGID  STAT  UID  TIME  COMMAND
0      2      0      0  ?      -1  S      0  0:00  [kthreadd]
2      3      0      0  ?      -1  S      0  0:00  \_ [ksoftirqd/0]
2      5      0      0  ?      -1  S      0  0:00  \_ [kworker/u:0]
2      6      0      0  ?      -1  S      0  0:00  \_ [migration/0]
2      7      0      0  ?      -1  S      0  0:00  \_ [watchdog/0]
2      8      0      0  ?      -1  S      0  0:00  \_ [migration/1]
2     10      0      0  ?      -1  S      0  0:00  \_ [ksoftirqd/1]
2     11      0      0  ?      -1  S      0  0:00  \_ [watchdog/1]
2     12      0      0  ?      -1  S<      0  0:00  \_ [cpuset]
2     13      0      0  ?      -1  S<      0  0:00  \_ [khelper]
2     14      0      0  ?      -1  S      0  0:00  \_ [kdevtmpfs]
2     15      0      0  ?      -1  S<      0  0:00  \_ [netns]
2     16      0      0  ?      -1  S      0  0:00  \_ [sync_supers]
2     17      0      0  ?      -1  S      0  0:00  \_ [bdi-default]
2     18      0      0  ?      -1  S<      0  0:00  \_ [kintegrityd]
2     19      0      0  ?      -1  S<      0  0:00  \_ [kblockd]
2     20      0      0  ?      -1  S<      0  0:00  \_ [ata_sff]
2     21      0      0  ?      -1  S      0  0:00  \_ [khubd]
2     22      0      0  ?      -1  S<      0  0:00  \_ [md]
2     23      0      0  ?      -1  S      0  0:00  \_ [kworker/1:1]
2     25      0      0  ?      -1  S      0  0:00  \_ [kswapd0]
2     26      0      0  ?      -1  SN      0  0:00  \_ [ksmd]
2     27      0      0  ?      -1  SN      0  0:00  \_ [khugepaged]
2     28      0      0  ?      -1  S      0  0:00  \_ [fsnotify_mark]
2     29      0      0  ?      -1  S<      0  0:00  \_ [crypto]
2     35      0      0  ?      -1  S<      0  0:00  \_ [kthrotld]
2     38      0      0  ?      -1  S      0  0:00  \_ [scsi_eh_0]
2     39      0      0  ?      -1  S      0  0:00  \_ [scsi_eh_1]
2     40      0      0  ?      -1  S      0  0:00  \_ [scsi_eh_2]
2     41      0      0  ?      -1  S      0  0:00  \_ [kworker/u:2]
2     43      0      0  ?      -1  S<      0  0:00  \_ [kpsmouse]
2     44      0      0  ?      -1  S<      0  0:00  \_ [deferwq]
2     46      0      0  ?      -1  S      0  0:00  \_ [kworker/0:2]
2    238      0      0  ?      -1  S      0  0:00  \_ [kworker/1:2]
2    290      0      0  ?      -1  S<      0  0:00  \_ [kdmflush]
2    291      0      0  ?      -1  S<      0  0:00  \_ [kdmflush]
2    338      0      0  ?      -1  S      0  0:00  \_ [jbd2/dm-1-8]
2    339      0      0  ?      -1  S<      0  0:00  \_ [ext4-dio-unwrit]
2    375      0      0  ?      -1  S      0  0:00  \_ [kauditd]
```

UNIX-LIKE FOR LIFE

Friday, July 31, 2009

Linux process state codes

Here are the different values that the s, stat and state output specifiers (header "STAT" or "S") will display to describe the state of a process.

D Uninterruptible sleep (usually IO)

R Running or runnable (on run queue)

S Interruptible sleep (waiting for an event to complete)

T Stopped, either by a job control signal or because it is being traced.

W paging (not valid since the 2.6.xx kernel)

X dead (should never be seen)

Z Defunct ("zombie") process, terminated but not reaped by its parent.

For BSD formats and when the stat keyword is used, additional characters may be displayed:

< high-priority (not nice to other users)

N low-priority (nice to other users)

L has pages locked into memory (for real-time and custom IO)

s is a session leader

l is multi-threaded (using CLONE_THREAD, like NPTL pthreads do)

+ is in the foreground process group

Posted by M.Burak Alkan at 12:01 AM

Labels [linux](#)

No comments:

[Post a Comment](#)

[Newer Post](#)

[Home](#)

[Old](#)

Subscribe to: [Post Comments \(Atom\)](#)





Process Control Block (PCB)

Information associated with each process
(also called **task control block**)

Process state – running, waiting, etc

Program counter – location of instruction to next execute


CPU registers – contents of all process-centric registers

CPU scheduling information- priorities, scheduling queue pointers

Memory-management information – memory allocated to the process


Accounting information – CPU used, clock time elapsed since start, time limits

I/O status information – I/O devices allocated to process, list of open files





CPU Switch From Process to Process





Process Control Block (PCB)

include/linux/sched.h

C/C++ - my_kernel/include/linux/sched.h - Eclipse Platform

File Edit Source Refactor Navigate Search Project Run Window Help

Project Explorer Outline

task_struct

- state : volatile long
- stack : void*
- usage : atomic_t
- flags : unsigned int
- ptrace : unsigned int
- wake_entry : struct llist_node
- on_cpu : int
- on_rq : int
- prio : int
- static_prio : int
- normal_prio : int
- rt_priority : unsigned int
- sched_class : const struct sched_class*
- se : struct sched_entity
- rt : struct sched_rt_entity
- preempt_notifiers : struct hlist_head
- fpu_counter : unsigned char
- btrace_seq : unsigned int
- policy : unsigned int
- nr_cpus_allowed : int
- cpus_allowed : cpumask_t
- rcu_read_lock_nesting : int
- rcu_read_unlock_special : char
- rcu_node_entry : struct list_head
- rcu_blocked_node : struct rcu_node*
- rcu_boost_mutex : struct rt_mutex*
- sched_info : struct sched_info
- tasks : struct list_head

```
struct task_struct {
    volatile long state; /* -1 unrunnable, 0 runnable, >0 stopped */
    void *stack;
    atomic_t usage;
    unsigned int flags; /* per process flags, defined below */
    unsigned int ptrace;

#ifdef CONFIG_SMP
    struct llist_node wake_entry;
    int on_cpu;
#endif
    int on_rq;

    int prio, static_prio, normal_prio;
    unsigned int rt_priority;
    const struct sched_class *sched_class;
    struct sched_entity se;
    struct sched_rt_entity rt;

#ifdef CONFIG_PREEMPT_NOTIFIERS
    /* list of struct preempt_notifier: */
    struct hlist_head preempt_notifiers;
#endif

    /*
     * fpu_counter contains the number of consecutive context switches
     * that the FPU is used. If this is over a threshold, the lazy fpu
     * saving becomes unlazy to save the trap. This is an unsigned char
     * so that after 256 times the counter wraps and the behavior turns
     * lazy again; this to deal with bursty apps that only use FPU for
     * a short time
     */
    unsigned char fpu_counter;
#ifdef CONFIG_BLK_DEV_IO_TRACE
    unsigned int btrace_seq;
#endif

    unsigned int policy;
    int nr_cpus_allowed;
    cpumask_t cpus_allowed;

#ifdef CONFIG_PREEMPT_RCU
    int rcu_read_lock_nesting;
    char rcu_read_unlock_special;
    struct list_head rcu_node_entry;
#endif /* #ifdef CONFIG_PREEMPT_RCU */
#ifdef CONFIG_TREE_PREEMPT_RCU
    struct rcu_node *rcu_blocked_node;
#endif /* #ifdef CONFIG_TREE_PREEMPT_RCU */
#ifdef CONFIG_RCU_BOOST
    struct rt_mutex *rcu_boost_mutex;
#endif /* #ifdef CONFIG_RCU_BOOST */
}
```

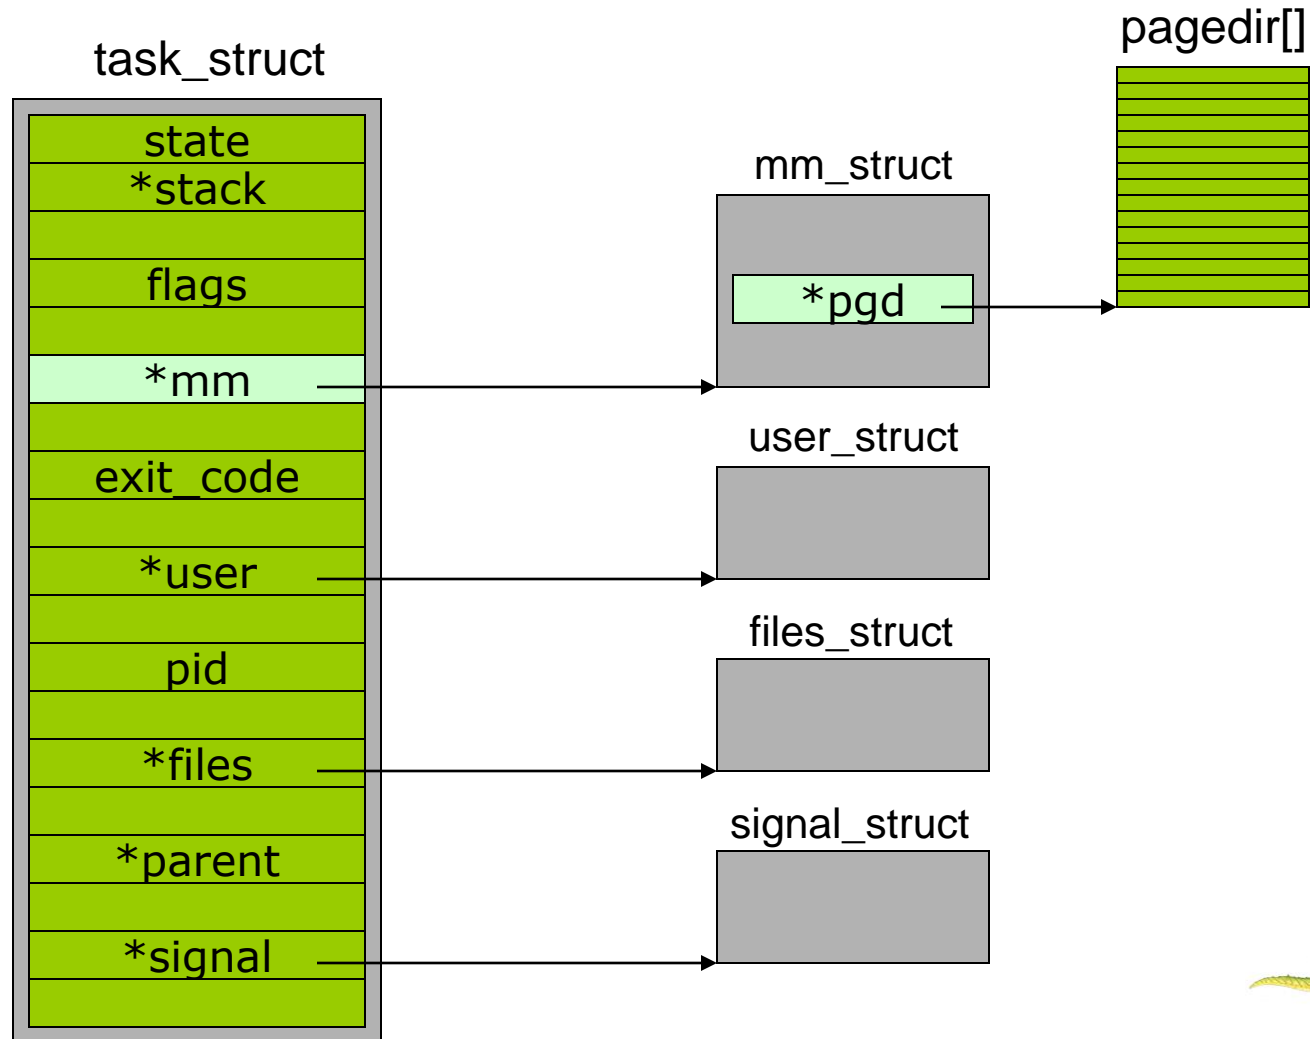


The Linux process descriptor

Each process descriptor contains many fields

and some are pointers to other kernel structures

which may themselves include fields that point to structures





Threads

So far, process has a single thread of execution

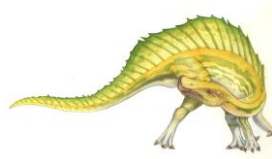
Consider having multiple program counters per process


Multiple locations can execute at once

- ▶ Multiple threads of control -> **threads**

Must then have storage for thread details, multiple program counters in PCB

See next chapter

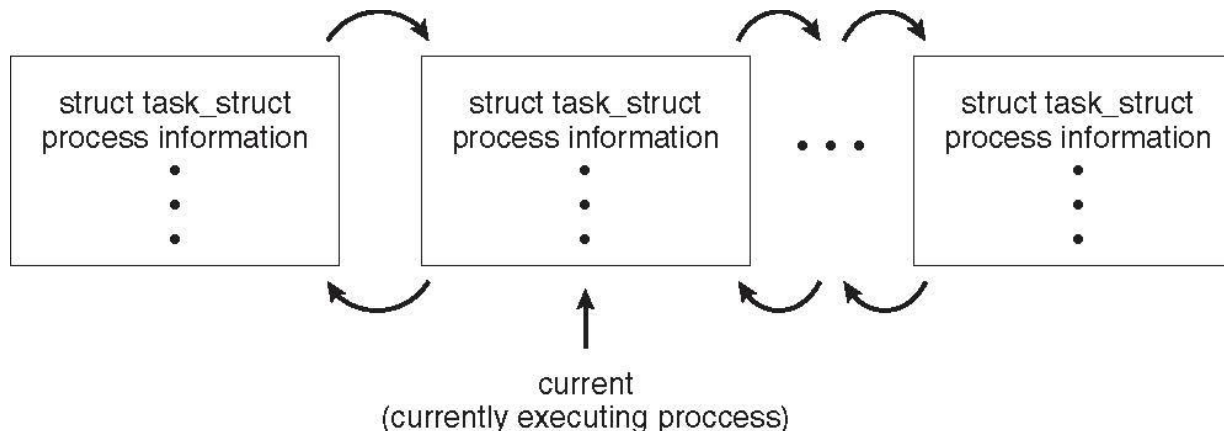




Process Representation in Linux

Represented by the C structure `task_struct`

```
pid t_pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```





Process Scheduling

Maximize CPU use, quickly switch processes onto CPU for time sharing

Process scheduler selects among available processes for next execution on CPU

Maintains **scheduling queues** of processes

Job queue – set of all processes in the system

Ready queue – set of all processes residing in main memory, ready and waiting to execute

Device queues – set of processes waiting for an I/O device

Processes migrate among the various queues








Representation of Process Scheduling

Queueing diagram represents queues, resources, flows





Schedulers

Short-term scheduler (or **CPU scheduler**) – selects which process should be executed next and allocates CPU

Sometimes the only scheduler in a system

Short-term scheduler is invoked frequently (milliseconds) \Rightarrow (must be fast)

Long-term scheduler (or **job scheduler**) – selects which processes should be brought into the ready queue

Long-term scheduler is invoked infrequently (seconds, minutes) \Rightarrow (may be slow)

The long-term scheduler controls the **degree of multiprogramming**

Processes can be described as either:

I/O-bound process – spends more time doing I/O than computations, many short CPU bursts

CPU-bound process – spends more time doing computations; few very long CPU bursts

Long-term scheduler strives for good ***process mix***






Addition of Medium Term Scheduling

Medium-term scheduler can be added if degree of multiple programming needs to decrease

Remove process from memory, store on disk, bring back in from disk to continue execution: **swapping**





Multitasking in Mobile Systems

Some mobile systems (e.g., early version of iOS) allow only one process to run, others suspended

Due to screen real estate, user interface limits iOS provides for a

Single **foreground** process- controlled via user interface

Multiple **background** processes– in memory, running, but not on the display, and with limits

Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback

Android runs foreground and background, with fewer limits

Background process uses a **service** to perform tasks

Service can keep running even if background process is suspended

Service has no user interface, small memory use





Context Switch

When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**

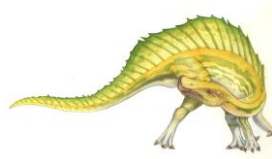
Context of a process represented in the PCB

Context-switch time is overhead; the system does no useful work while switching

The more complex the OS and the PCB → the longer the context switch

Time dependent on hardware support

Some hardware provides multiple sets of registers per CPU
→ multiple contexts loaded at once





Operations on Processes

System must provide mechanisms for:

- process creation,
- process termination,
- and so on as detailed next





Process Creation

Parent process create **children** processes, which, in turn create other processes, forming a **tree** of processes

Generally, process identified and managed via a **process identifier (pid)**

Resource sharing options

- Parent and children share all resources

- Children share subset of parent's resources

- Parent and child share no resources

Execution options


- Parent and children execute concurrently

- Parent waits until children terminate





A Tree of Processes in Linux





Process Creation (Cont.)

Address space


- Child duplicate of parent

- Child has a program loaded into it

UNIX examples

- fork()** system call creates new process

- exec()** system call used after a **fork()** to replace the process' memory space with a new program





C Program Forking Separate Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```





```
int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/lis", "lis", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait (NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

parent process memory

```
1230 int
1231 main(int argc, char **argv)
1232 {
1233     pid_t pid;
1234     /* fork another process */
1235     struct pending *thispend;
1236     pid = fork();
1237     if (pid < 0) { /* error occurred */
1238         /* Signals that are required the number of such signals. */
1239         static int const sig1 =
1240         {
1241             /* This one is handled specially. */
1242             SIGTSTP,
1243         };
1244         /* The usual suspects */
1245         SIGPROF, SIGPOLL, SIGRTM, SIGTALM, SIGTDM, SIGTSTP, SIGTTRM,
1246         #ifdef SIGPOLL
1247             SIGPOLL
1248         #endif
1249         #ifdef SIGPROF
1250             SIGPROF
1251         #endif
1252         #ifdef SIGTALM
1253             SIGTALM,
1254         #endif
1255         #ifdef SIGXCPU
1256             SIGXCPU
1257         #endif
1258         #ifdef SIGXFSZ
1259             SIGXFSZ,
1260         #endif
1261     };
1262     enum { nsigs = ARRAY_CARDINALITY (sig) };
1263 }
```

child process memory





Creating a Separate Process via Windows API

```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```





Process Termination

Process executes last statement and then asks the operating system to delete it using the **exit()** system call.

Returns status data from child to parent (via **wait()**)

Process' resources are deallocated by operating system

Parent may terminate the execution of children processes using the **abort()** system call. Some reasons for doing so:

Child has exceeded allocated resources

Task assigned to child is no longer required

The parent is exiting and the operating systems does not allow a child to continue if its parent terminates





Process Termination

Some operating systems do not allow child to exist if its parent has terminated. If a process terminates, then all its children must also be terminated.

cascading termination. All children, grandchildren, etc. are terminated.

The termination is initiated by the operating system.

The parent process may wait for termination of a child process by using the `wait()` system call. The call returns status information and the pid of the terminated process

```
pid = wait(&status);
```

If no parent waiting (did not invoke `wait()`) process is a **zombie**

If parent terminated without invoking `wait`, process is an **orphan**





Multiprocess Architecture – Chrome Browser

Many web browsers ran as single process (some still do)

If one web site causes trouble, entire browser can hang or crash

Google Chrome Browser is multiprocess with 3 different types of processes:

Browser process manages user interface, disk and network I/O

Renderer process renders web pages, deals with HTML, Javascript. A new renderer created for each website opened

- ▶ Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits

Plug-in process for each type of plug-in



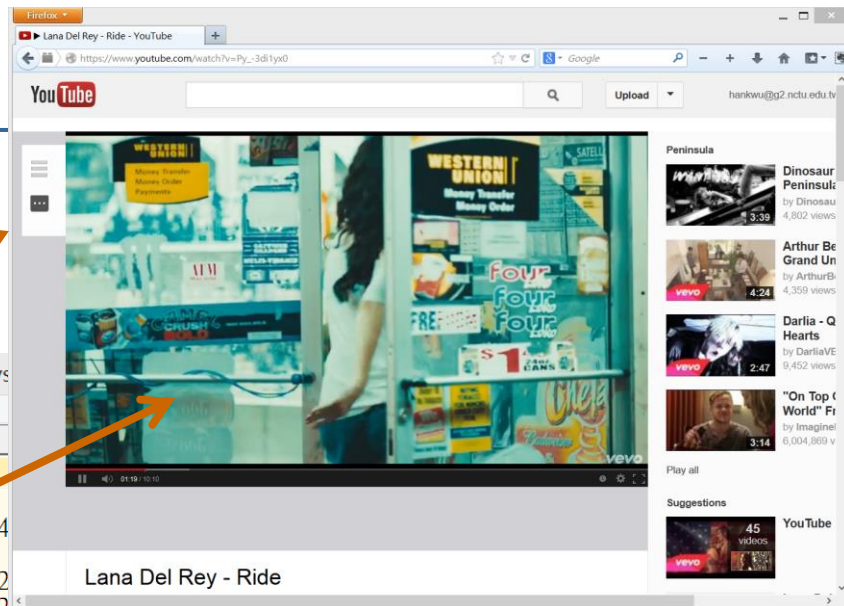


Process Explorer - Sysinternals: www.sys

Process	PID	CPU
putty.exe	11228	
POWERPNT.EXE	5512	4
procexp.exe	7760	
procexp64.exe	8224	2.092
firefox.exe	3564	0.412
plugin-container.exe	5404	0.34
FlashPlayerPlugin_11_8_800_168.exe	9140	4.652
FlashPlayerPlugin_11_8_800_168.exe	8848	2.58
jusched.exe	8888	2.480
iTunesHelper.exe	7404	< 0.01
pageant.exe	8492	1.044
AdobeARM.exe	2608	0.01
chrome.exe	7768	0.02
chrome.exe	2136	131.6
chrome.exe	7500	94.724

Type	Name
Key	HKLM\SYSTEM\ControlSet001\Services\WinSock2\Parameters\NameSpace_0
Key	HKLM\SOFTWARE\Microsoft\Windows\CurrentVersion\MMDevices\Audio\Re
Key	HKLM\SOFTWARE\Microsoft\Windows\CurrentVersion\MMDevices\Audio\Re
Key	HKCU\Software\Classes
Mutant	\Sessions\8\BaseNamedObjects\MSCTF.Asm.MutexDefault8
Section	\Windows\Theme3398143603
Section	\Sessions\8\Windows\Theme3468387045
Section	\BaseNamedObjects__ComCatalogCache__
Section	\BaseNamedObjects__ComCatalogCache__
Section	...\ASqmManifestVersion
Section	...\ASqmManifest_27eb
Section	\Sessions\8\BaseNamedObjects\windows_shell_global_counters
Thread	FlashPlayerPlugin_11_8_800_168.exe(8848): 4504
Thread	FlashPlayerPlugin_11_8_800_168.exe(8848): 10116
Thread	FlashPlayerPlugin_11_8_800_168.exe(8848): 4504

CPU Usage: 8.17% Commit Charge: 29.74% Processes: 95 Physical Usage: 29.95%



Details 安全性

Basic Information

Name: \Windows\Theme3398143603

Type: Section

Description: A memory mapped file or paging-file backed memory region.

Address: 0xFFFFF8A003B62E60

References

References: 13893608

Handles: 53

Quota Charges

Paged: 2272

Non-Paged: 168

Section Info

Type: COMMITED

Size: 0x00115000

確定 取消



Interprocess Communication

Processes within a system may be ***independent*** or ***cooperating***

Cooperating process can affect or be affected by other processes, including sharing data

Reasons for cooperating processes:

- Information sharing

- Computation speedup

- Modularity

- Convenience

Cooperating processes need **interprocess communication (IPC)**

Two models of IPC

- Shared memory**

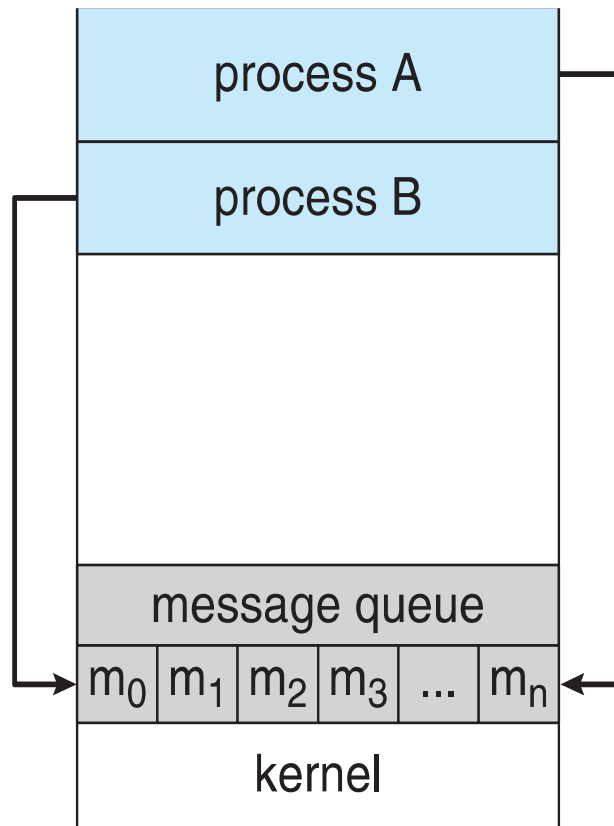
- Message passing**






Communications Models

(a) Message passing. (b) shared memory.



(a)



(b)





Cooperating Processes

Independent process cannot affect or be affected by the execution of another process

Cooperating process can affect or be affected by the execution of another process

Advantages of process cooperation

- Information sharing

- Computation speed-up

- Modularity

- Convenience





Producer-Consumer Problem

Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process

unbounded-buffer places no practical limit on the size of the buffer

bounded-buffer assumes that there is a fixed buffer size





Bounded-Buffer – Shared-Memory Solution


Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

Solution is correct, but can only use $\text{BUFFER_SIZE}-1$ elements





Bounded-Buffer – Producer

```
item next_produced;
while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```





Bounded Buffer – Consumer

```
item next_consumed;
while (true) {
    while (in == out)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next consumed */
}
```





Interprocess Communication – Shared Memory

An area of memory shared among the processes that wish to communicate

The communication is under the control of the users processes not the operating system.

Major issues is to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory.

Synchronization is discussed in great details in Chapter 5.





Interprocess Communication – Message Passing

Mechanism for processes to communicate and to synchronize their actions

Message system – processes communicate with each other without resorting to shared variables

IPC facility provides two operations:

send(*message*)

receive(*message*)

The *message* size is either fixed or variable





Message Passing (Cont.)

If processes P and Q wish to communicate, they need to:

- Establish a **communication link** between them

- Exchange messages via send/receive

Implementation issues:

- How are links established?

- Can a link be associated with more than two processes?

- How many links can there be between every pair of communicating processes?

- What is the capacity of a link?

- Is the size of a message that the link can accommodate fixed or variable?

- Is a link unidirectional or bi-directional?





Message Passing (Cont.)

Implementation of communication link

Physical:

- ▶ Shared memory
- ▶ Hardware bus
- ▶ Network

Logical:

- ▶ Direct or indirect
- ▶ Synchronous or asynchronous
- ▶ Automatic or explicit buffering





Direct Communication

Processes must name each other explicitly:

send (P , *message*) – send a message to process P

receive(Q , *message*) – receive a message from process Q

Properties of communication link

Links are established automatically

A link is associated with exactly one pair of communicating processes

Between each pair there exists exactly one link

The link may be unidirectional, but is usually bi-directional





Indirect Communication

Messages are directed and received from mailboxes (also referred to as ports)

- Each mailbox has a unique id

- Processes can communicate only if they share a mailbox

Properties of communication link

- Link established only if processes share a common mailbox

- A link may be associated with many processes

- Each pair of processes may share several communication links

- Link may be unidirectional or bi-directional





Indirect Communication

Operations

- create a new mailbox (port)

- send and receive messages through mailbox

- destroy a mailbox

Primitives are defined as:

send(*A*, *message*) – send a message to mailbox *A*

receive(*A*, *message*) – receive a message from mailbox *A*





Indirect Communication

Mailbox sharing

P_1 , P_2 , and P_3 share mailbox A

P_1 , sends; P_2 and P_3 receive

Who gets the message?

Solutions

Allow a link to be associated with at most two processes

Allow only one process at a time to execute a receive operation

Allow the system to select arbitrarily the receiver.

Sender is notified who the receiver was.





Synchronization

Message passing may be either blocking or non-blocking

Blocking is considered **synchronous**

Blocking send -- the sender is blocked until the message is received

Blocking receive -- the receiver is blocked until a message is available

Non-blocking is considered **asynchronous**

Non-blocking send -- the sender sends the message and continue

Non-blocking receive -- the receiver receives:

A valid message, or

Null message

Different combinations possible

If both send and receive are blocking, we have a **rendezvous**





Synchronization (Cont.)

Producer-consumer becomes trivial

```
message next_produced;  
while (true) {  
    /* produce an item in next produced */  
    send(next_produced);  
}
```

```
message next_consumed;  
while (true) {  
    receive(next_consumed);  
  
    /* consume the item in next consumed */  
}
```





Buffering

Queue of messages attached to the link.

implemented in one of three ways

1. Zero capacity – no messages are queued on a link.
Sender must wait for receiver (rendezvous)
2. Bounded capacity – finite length of n messages
Sender must wait if link full
3. Unbounded capacity – infinite length
Sender never waits





Examples of IPC Systems - POSIX

POSIX Shared Memory

Process first creates shared memory segment

```
shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);
```

Also used to open an existing segment to share it

Set the size of the object

```
ftruncate(shm_fd, 4096);
```

Now the process could write to the shared memory

```
sprintf(shared_memory, "Writing to shared  
memory");
```





IPC POSIX Producer

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr,"%s",message_0);
    ptr += strlen(message_0);
    sprintf(ptr,"%s",message_1);
    ptr += strlen(message_1);

    return 0;
}
```





IPC POSIX Consumer

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```





Examples of IPC Systems - Mach

Mach communication is message based

Even system calls are messages

Each task gets two mailboxes at creation- Kernel and Notify

Only three system calls needed for message transfer

`msg_send()` , `msg_receive()` , `msg_rpc()`

Mailboxes needed for communication, created via

`port_allocate()`

Send and receive are flexible, for example four options if mailbox full:

- ▶ Wait indefinitely
- ▶ Wait at most n milliseconds
- ▶ Return immediately
- ▶ Temporarily cache a message





Examples of IPC Systems – Windows

Message-passing centric via **advanced local procedure call (LPC)** facility

Only works between processes on the same system

Uses ports (like mailboxes) to establish and maintain communication channels


Communication works as follows:

- ▶ The client opens a handle to the subsystem's **connection port** object.
- ▶ The client sends a connection request.
- ▶ The server creates two private **communication ports** and returns the handle to one of them to the client.
- ▶ The client and server use the corresponding port handle to send messages or callbacks and to listen for replies.





Local Procedure Calls in Windows





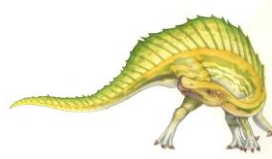
Communications in Client-Server Systems

Sockets

Remote Procedure Calls

Pipes

Remote Method Invocation (Java)





Sockets

A **socket** is defined as an endpoint for communication

Concatenation of IP address and **port** – a number included at start of message packet to differentiate network services on a host

The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**

Communication consists between a pair of sockets


All ports below 1024 are **well known**, used for standard services

Special IP address 127.0.0.1 (**loopback**) to refer to system on which process is running





Socket Communication





Sockets in Java

Three types of sockets

**Connection-oriented
(TCP)**

Connectionless (UDP)

MulticastSocket
class— data can be sent
to multiple recipients

Consider this “Date” server:

```
import java.net.*;
import java.io.*;

public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

            /* now listen for connections */
            while (true) {
                Socket client = sock.accept();

                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());

                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```





Remote Procedure Calls

Remote procedure call (RPC) abstracts procedure calls between processes on networked systems

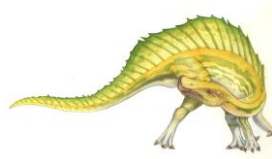
Again uses ports for service differentiation

Stubs – client-side proxy for the actual procedure on the server

The client-side stub locates the server and **marshalls** the parameters

The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server

On Windows, stub code compile from specification written in **Microsoft Interface Definition Language (MIDL)**





Remote Procedure Calls (Cont.)

Data representation handled via **External Data Representation (XDL)** format to account for different architectures

Big-endian and **little-endian**

Remote communication has more failure scenarios than local

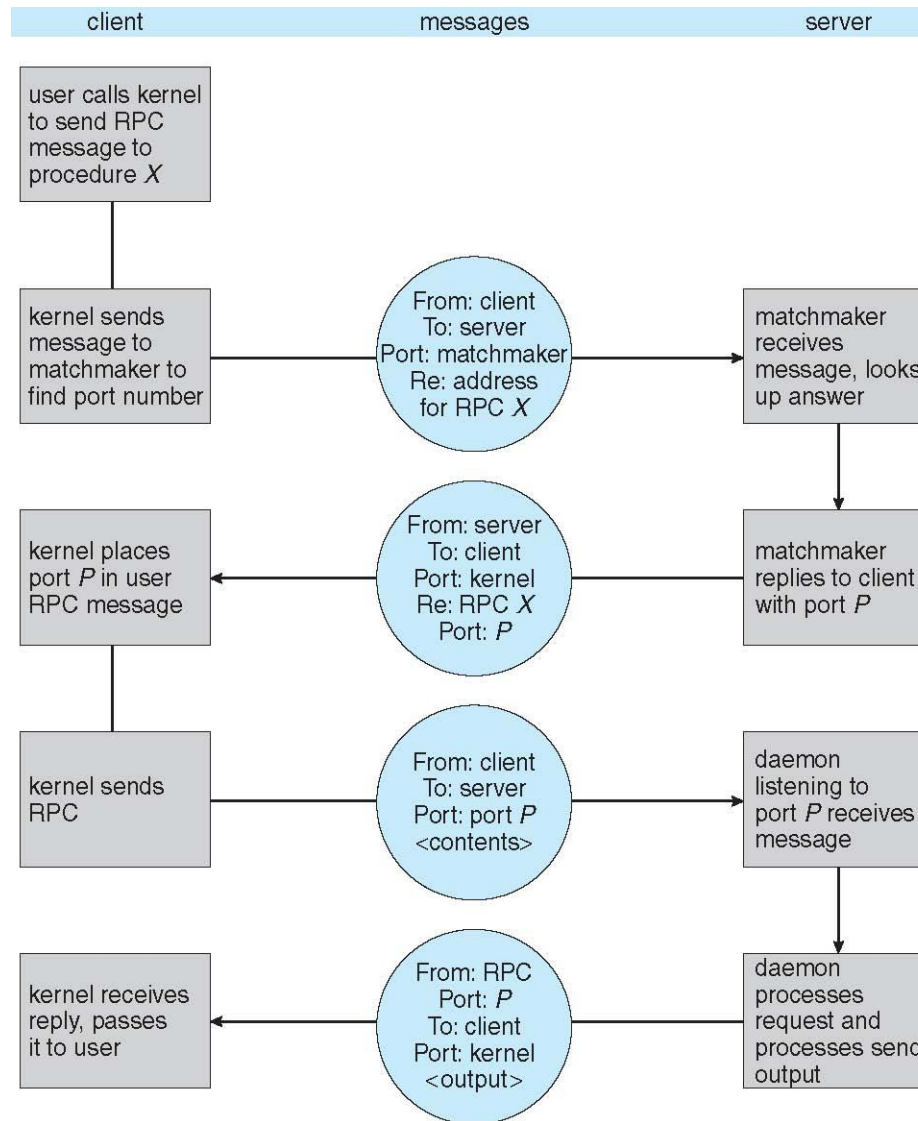
Messages can be delivered ***exactly once*** rather than ***at most once***

OS typically provides a rendezvous (or **matchmaker**) service to connect client and server





Execution of RPC





Pipes

Acts as a conduit allowing two processes to communicate

Issues:

Is communication unidirectional or bidirectional?

In the case of two-way communication, is it half or full-duplex?

Must there exist a relationship (i.e., ***parent-child***) between the communicating processes?

Can the pipes be used over a network?

Ordinary pipes – cannot be accessed from outside the process that created it. Typically, a parent process creates a pipe and uses it to communicate with a child process that it created.

Named pipes – can be accessed without a parent-child relationship.





Ordinary Pipes


Ordinary Pipes allow communication in standard producer-consumer style

Producer writes to one end (the **write-end** of the pipe)

Consumer reads from the other end (the **read-end** of the pipe)

Ordinary pipes are therefore unidirectional

Require parent-child relationship between communicating processes



Windows calls these **anonymous pipes**

See Unix and Windows code samples in textbook





Named Pipes

Named Pipes are more powerful than ordinary pipes

Communication is bidirectional

No parent-child relationship is necessary between the communicating processes

Several processes can use the named pipe for communication

Provided on both UNIX and Windows systems



End of Chapter 3

