

Chapter 5: Process Synchronization





Chapter 5: Process Synchronization

- Background
- The Critical-Section Problem
- Peterson's Solution
- Synchronization Hardware
- Mutex Locks
- Semaphores
- Classic Problems of Synchronization
- Monitors
- Synchronization Examples
- Alternative Approaches





Objectives

To present the concept of process synchronization.

To introduce the critical-section problem, whose solutions can be used to ensure the consistency of shared data

To present both software and hardware solutions of the critical-section problem

To examine several classical process-synchronization problems

To explore several tools that are used to solve process synchronization problems





Background

Processes can execute concurrently

May be interrupted at any time, partially completing execution

Concurrent access to shared data may result in data inconsistency

Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes

Illustration of the problem:

Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers. We can do so by having an integer **counter** that keeps track of the number of full buffers. Initially, **counter** is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.







Synchronization

```
hank@Maestro:~/git_Hank/courses/2012_os/code/synchronization
File Edit View Search Terminal Help
top - 23:31:11 up 21 min, 2 users, load average: 0.00, 0.04, 0.10
Tasks: 140 total, 1 running, 139 sleeping, 0 stopped, 0 zombie
Cpu0 : 0.3%us, 2.0%sy, 0.0%ni, 97.6%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu1 : 0.0%us, 0.3%sy, 0.0%ni, 99.7%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Cpu2 : 0.3%us, 0.7%sy, 0.0%ni, 99.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 1019476k total, 933240k used, 86236k free, 37060k buffers
Swap: 2064380k total, 0k used, 2064380k free, 474464k cached
```

Type '1' to toggle SMP view

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
796	root	20	0	299m	79m	11m	S	4.3	8.0	0:43.49	Xorg
1277	hank	20	0	1764m	100m	43m	S	4.0	10.1	0:37.64	gnome-shell
1497	hank	20	0	572m	15m	9892	S	0.3	1.6	0:01.51	gnome-terminal
1636	hank	20	0	15256	1216	900	R	0.3	0.1	0:00.15	top
1	root	20	0	67280	25m	2080	S	0.0	2.5	0:01.25	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kthreadd
3	root	20	0	0	0	0	S	0.0	0.0	0:00.11	ksoftirqd
5	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	kworker/
7	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	kworker/
8	root	RT	0	0	0	0	S	0.0	0.0	0:00.00	migratio
9	root	RT	0	0	0	0	S	0.0	0.0	0:00.02	watchdog
10	root	RT	0	0	0	0	S	0.0	0.0	0:00.36	migratio
12	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	kworker/
13	root	20	0	0	0	0	S	0.0	0.0	0:00.06	ksoftirqd
14	root	RT	0	0	0	0	S	0.0	0.0	0:00.01	watchdog
15	root	RT	0	0	0	0	S	0.0	0.0	0:00.01	migratio
17	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	kworker/
18	root	20	0	0	0	0	S	0.0	0.0	0:00.02	ksoftirqd
19	root	RT	0	0	0	0	S	0.0	0.0	0:00.01	watchdog
20	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	cpuset
21	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	khelper
22	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kdevtmpf
23	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	netns
24	root	20	0	0	0	0	S	0.0	0.0	0:00.00	bdi-defa
25	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	kintegri

Make sure your machine has multiple cores





Single thread addition

```
unsigned long long int cnt;

void* AdditionX(void* pParam);
void* AdditionY(void* pParam);

int main() {

    timespec start_realtime, end_realtime;
    const unsigned long long int limit = 1000000000;
    cnt = 0;

    clock_gettime(CLOCK_REALTIME, &start_realtime );
    printf("start running...\n");

    AdditionX((void*)&limit);
    AdditionY((void*)&limit);

    printf("cnt = %llu\n", cnt);

    clock_gettime(CLOCK_REALTIME, &end_realtime );

    printf("duration = %d nanoseconds\n", (end_realtime.tv_sec - start_realtime.tv_sec)*1000000000 +
        (end_realtime.tv_nsec - start_realtime.tv_nsec));

    return 0;
}
```





Single thread addition

```
void* AdditionX(void* pParam)
{
    unsigned long long int limit = *((const unsigned long long int*)pParam);

    printf("thread X limit = %d\n", limit);

    while(limit-- > 0) {
        cnt++;
    }

    return 0;
}
```

```
void* AdditionY(void* pParam)
{
    unsigned long long int limit = *((const unsigned long long int*)pParam);

    printf("thread Y limit = %d\n", limit);

    while(limit-- > 0 ) {
        cnt++;
    }

    return 0;
}
```





Single thread addition

```
hank@Maestro:~/git_Hank/courses/2012_os/code/synchronization

File Edit View Search Terminal Help
remote: Counting objects: 11, done.
remote: Compressing objects: 100% (4/4), done.
remote: Total 6 (delta 2), reused 0 (delta 0)
Jnpacking objects: 100% (6/6), done.
From ssh://code.cs.nctu.edu.tw/var/git/person/Hank
    dbdda1..464d0c1 master    -> origin/master
Jpdating dbdda1..464d0c1
Fast-forward
  courses/2012_os/code/synchronization/complex.cpp | 80 -----
  1 file changed, 80 deletions(-)
  delete mode 100644 courses/2012_os/code/synchronization/complex.cpp
[hank@Maestro synchronization]$ ls -al
total 20
drwxrwxr-x. 2 hank hank 4096 Oct 31 23:45 .
drwxrwxr-x. 6 hank hank 4096 Oct 31 21:21 ..
-rw-rw-r--. 1 hank hank 1399 Oct 31 22:27 multiple_thread_add.cpp
-rw-rw-r--. 1 hank hank 1502 Oct 31 22:30 multiple_thread_inc.cpp
-rw-rw-r--. 1 hank hank 1100 Oct 31 21:47 single_thread_add.cpp
[hank@Maestro synchronization]$
[hank@Maestro synchronization]$ g++ -g ./single_thread_add.cpp -lrt -lpthread
[hank@Maestro synchronization]$ ./a.out
start running...
thread X limit = 100000000
thread Y limit = 100000000
cnt = 200000000
duration = 456750331 nanoseconds
[hank@Maestro synchronization]$
```





Multithread addition

```
unsigned long long int cnt;

void* AdditionX(void* pParam);
void* AdditionY(void* pParam);

int main() {

    pthread_t tidX, tidY;
    pthread_attr_t thread_attr;
    timespec start_realtime, end_realtime;

    const unsigned long long int limit = 100000000;

    cnt = 0;

    pthread_attr_init(&thread_attr);
    pthread_attr_setdetachstate(&thread_attr, PTHREAD_CREATE_JOINABLE);

    clock_gettime(CLOCK_REALTIME, &start_realtime );
    printf("start running...\n");

    pthread_create(&tidX, &thread_attr, AdditionX, (void*)&limit );
    pthread_create(&tidY, &thread_attr, AdditionY, (void*)&limit );

    pthread_join(tidX,0);
    pthread_join(tidY,0);

    printf("cnt = %llu\n", cnt);

    clock_gettime(CLOCK_REALTIME, &end_realtime );

    printf("duration = %d nanoseconds\n", (end_realtime.tv_sec - start_realtime.tv_sec)*1000000000 +
        (end_realtime.tv_nsec - start_realtime.tv_nsec));

    return 0;
}
```





Multithread Addition

```
[hank@Maestro synchronization]$  
[hank@Maestro synchronization]$ g++ -g ./multiple_thread_add.cpp -lrt -lpthread  
[hank@Maestro synchronization]$ ./a.out  
start running...  
thread X limit = 1000000000  
thread Y limit = 1000000000  
cnt = 112770930  
duration = 733176733 nanoseconds  
[hank@Maestro synchronization]$
```





Multithread Addition

```
void* AdditionX(void* pParam)
{
    unsigned long long int limit =
        *((const unsigned long long int*)pParam);

    printf("thread X limit = %d\n", limit);

    while(limit-- > 0) {
        cnt++;
    }

    return 0;
}
```

```
void* AdditionY(void* pParam)
{
    unsigned long long int limit =
        *((const unsigned long long int*)pParam);

    printf("thread Y limit = %d\n", limit);

    while(limit-- > 0 ) {
        cnt++;
    }

    return 0;
}
```

The code of AdditionX and the code of AdditionY may interleave

Any problem?

The code of AdditionX and the code of AdditionY may overlap

Any problem?





Multithread Addition

```
for function AdditionX(void*):
<+0>:    push    %rbp
<+1>:    mov     %rsp,%rbp
<+4>:    sub     $0x20,%rsp
<+8>:    mov     %rdi,-0x18(%rbp)
<+12>:   mov     -0x18(%rbp),%rax
<+16>:   mov     (%rax),%rax
<+19>:   mov     %rax,-0x8(%rbp)
<+23>:   mov     -0x8(%rbp),%rax
<+27>:   mov     %rax,%rsi
<+30>:   mov     $0x400bb8,%edi
<+35>:   mov     $0x0,%eax
<+40>:   callq   0x400740 <printf@plt>
<+45>:   jmp     0x400a59 <AdditionX(void*)+65>
<+47>:   mov     0x2005aa(%rip),%rax      # 0x600ff8 <cnt>
<+54>:   add     $0x1,%rax
<+58>:   mov     %rax,0x20059f(%rip)     # 0x600ff8 <cnt>
<+65>:   cmpq    $0x0,-0x8(%rbp)
<+70>:   setne    %al
<+73>:   subq    $0x1,-0x8(%rbp)
<+78>:   test    %al,%al
<+80>:   jne     0x400a47 <AdditionX(void*)+47>
<+82>:   mov     $0x0,%eax
<+87>:   leaveq   %eax
<+88>:   retq
```

```
for function AdditionY(void*):
<+0>:    push    %rbp
<+1>:    mov     %rsp,%rbp
<+4>:    sub     $0x20,%rsp
<+8>:    mov     %rdi,-0x18(%rbp)
<+12>:   mov     -0x18(%rbp),%rax
<+16>:   mov     (%rax),%rax
<+19>:   mov     %rax,-0x8(%rbp)
<+23>:   mov     -0x8(%rbp),%rax
<+27>:   mov     %rax,%rsi
<+30>:   mov     $0x400bce,%edi
<+35>:   mov     $0x0,%eax
<+40>:   callq   0x400740 <printf@plt>
<+45>:   jmp     0x400ab2 <AdditionY(void*)+65>
<+47>:   mov     0x200551(%rip),%rax      # 0x600ff8 <cnt>
<+54>:   add     $0x1,%rax
<+58>:   mov     %rax,0x200546(%rip)     # 0x600ff8 <cnt>
<+65>:   cmpq    $0x0,-0x8(%rbp)
<+70>:   setne    %al
<+73>:   subq    $0x1,-0x8(%rbp)
<+78>:   test    %al,%al
<+80>:   jne     0x400aa0 <AdditionY(void*)+47>
<+82>:   mov     $0x0,%eax
<+87>:   leaveq   %eax
<+88>:   retq
```

'cnt++' is compiled into multiple instructions





Multithread Inc

```
void* AdditionX(void* pParam)
{
    unsigned long long int limit = *((const unsigned long long int*)pParam);

    printf("thread X limit = %d\n", limit);

    while(limit-- > 0) {
        asm ("incq %0": "=m" (cnt)); ←
    }

    return 0;
}
```

A single line of assembly

Will this version run correctly on a uniprocessor?


```
void* AdditionY(void* pParam)
{
    unsigned long long int limit = *((const unsigned long long int*)pParam);

    printf("thread Y limit = %d\n", limit);

    while(limit-- > 0 ) {
        asm ("incq %0": "=m" (cnt));
    }

    return 0;
}
```

Will this version run correctly on a multiprocessor?





Producer

```
while (true) {  
    /* produce an item in next produced */  
  
    while (counter == BUFFER_SIZE) ;  
        /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```





Consumer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```





Race Condition

`counter++` could be implemented as

```
register1 = counter
register1 = register1 + 1
counter = register1
```

`counter--` could be implemented as

```
register2 = counter
register2 = register2 - 1
counter = register2
```

Consider this execution interleaving with “count = 5” initially:

S0: producer execute	<code>register1 = counter</code>	{register1 = 5}
S1: producer execute	<code>register1 = register1 + 1</code>	{register1 = 6}
S2: consumer execute	<code>register2 = counter</code>	{register2 = 5}
S3: consumer execute	<code>register2 = register2 - 1</code>	{register2 = 4}
S4: producer execute	<code>counter = register1</code>	{counter = 6}
S5: consumer execute	<code>counter = register2</code>	{counter = 4}





Critical Section

Producer thread

```
while (true) {  
  
    /* produce an item and put in nextProduced */  
  
    while (count == BUFFER_SIZE); // do nothing  
    buffer[in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```

Consumer thread

```
while (true) {  
    while (count == 0) ; // do nothing  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--; /* consume the item in nextConsumed */  
}
```

Code in this "critical section" should not be interleaved or overlapped





Critical Section - How

CriticalSection cs;

Producer thread

```
while (true) {  
  
    /* produce an item and put in nextProduced  
    */  
  
    while (count == BUFFER_SIZE ); // do  
    nothing  
  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    cs.lock();  
    count++;  
    cs.unlock();  
}
```

Consumer thread

```
while (true) {  
    while (count == 0) ; // do nothing  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    cs.lock();  
    count--; /* consume the item in nextConsumed  
    cs.unlock();  
}
```





Critical Section Problem

Consider system of n processes $\{p_0, p_1, \dots, p_{n-1}\}$

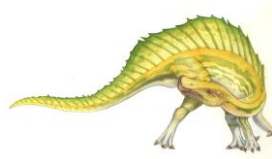
Each process has **critical section** segment of code

Process may be changing common variables, updating table, writing file, etc

When one process in critical section, no other may be in its critical section

Critical section problem is to design protocol to solve this

Each process must ask permission to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**





Critical Section

General structure of process P_i

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true);
```





Algorithm for Process P_i

```
do {  
    while (turn == j);  
        critical section  
    turn = j;  
        remainder section  
} while (true);
```





Solution to Critical-Section Problem

1. **Mutual Exclusion** - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections
2. **Progress (Freedom from Deadlock)** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely
3. **Bounded Waiting (Freedom from Starvation)** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted
 - Assume that each process executes at a nonzero speed
 - No assumption concerning **relative speed** of the n processes

“Freedom from starvation” implies “freedom from deadlock”





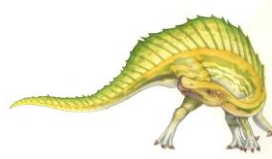
Critical-Section Handling in OS

Two approaches depending on if kernel is preemptive or non-preemptive

Preemptive – allows preemption of process when running in kernel mode

Non-preemptive – runs until exits kernel mode, blocks, or voluntarily yields CPU

- ▶ Essentially free of race conditions in kernel mode





1st attempt on the design of a Critical Section Lock

```
class CriticalSectionLockOne
{
    bool flag[2];

public:
    void lock()
    {
        int i = ThreadID.get();
        int j = 1-i;
        flag[i] = true;
        while(flag[j]) ; // wait
    }

    void unlock()
    {
        int i = ThreadID.get();
        flag[i] = false;
    }
};
```

Assume that the LOAD and STORE instructions are atomic; that is, cannot be interrupted.





1st attempt on the design of a Critical Section Lock

Does CriticalSectionLockOne satisfy
mutual exclusion?
freedom from deadlock?
freedom from starvation?





1st attempt on the design of a Critical Section Lock

Does CriticalSectionLockOne satisfy
mutual exclusion?
freedom from deadlock?
freedom from starvation?





1st attempt on the design of a Critical Section Lock

Proof of that CriticalSectionLockOne satisfies mutual exclusion

Assume thread A and B .

Let CS_A^j be the interval during which A executes the critical section for the j -th time.

Let CS_B^k be the interval during which B executes the critical section for the k -th time.

$\text{write}_A(x=v)$ denotes the event in which A assigns value v to field x .

$\text{read}_A(v==x)$ denotes the event in which A reads v from field x .

Suppose **CriticalSectionLockOne** does not satisfy mutual exclusion. Then there exist integers j and k such that $CS_A^j \rightarrow CS_B^k$ and $CS_B^k \rightarrow CS_A^j$. Consider each thread's last execution of the `lock()` method before entering its k -th (j -th) critical section.

Inspecting the code, we see that

$$\text{write}_A(\text{flag}[A]=\text{true}) \rightarrow \text{read}_A(\text{flag}[B]==\text{false}) \rightarrow CS_A \quad (1)$$

$$\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow \text{read}_B(\text{flag}[A]==\text{false}) \rightarrow CS_B \quad (2)$$


$$\text{read}_A(\text{flag}[B]==\text{false}) \rightarrow \text{write}_B(\text{flag}[B]=\text{true}) \quad (3)$$

By transitivity of the precedence order, we have

$$\begin{aligned} &\text{write}_A(\text{flag}[A]=\text{true}) \rightarrow \text{read}_A(\text{flag}[B]==\text{false}) \rightarrow \\ &\quad \text{write}_B(\text{flag}[B]=\text{true}) \rightarrow \text{read}_B(\text{flag}[A]==\text{false}) \end{aligned}$$

It follows that $\text{write}_A(\text{flag}[A]=\text{true}) \rightarrow \text{read}_B(\text{flag}[A]==\text{false})$ without an intervening write to the `flag[]` array, a contradiction. □






1st attempt on the design of a Critical Section Lock

CriticalSectionLockOne deadlocks if $\text{write}_A(\text{flag}[A]=\text{true})$ and $\text{write}_B(\text{flag}[B]=\text{true})$ events occur before $\text{read}_A(\text{flag}[B])$ and $\text{read}_B(\text{flag}[A])$ events, then both threads wait forever.





2nd attempt on the design of a Critical Section Lock

```
class CriticalSectionLockTwo
{
    int victim;

public:
    void lock()
    {
        int i = ThreadID.get();
        victim = i; // let the other go first
        while(victim == i) ; // wait
    }

    void unlock() {}
};
```

Exercise: prove that CriticalSectionLockTwo satisfies mutual exclusion.

But CriticalSectionLockTwo deadlocks if one thread runs completely before the other





Peterson's Solution

Two process solution

Assume that the LOAD and STORE instructions are atomic; that is, cannot be interrupted.

The two processes share two variables:

int **victim**;

Boolean **flag**[2]

The variable **victim** indicates whose turn it is to wait.

The **flag** array is used to indicate if a process is ready to enter the critical section. **flag[i]** = true implies that process P_i is ready!





Peterson Lock

```
class CriticalSection_PetersonLock
{
    bool flag[2];
    int victim;
public:
    void lock()
    {
        int i = ThreadID.get();
        int j = 1-i;
        flag[i] = true; // I'm ready
        victim = i; // you go first
        while(flag[j] && victim==i) ; // wait
    }

    void unlock()
    {
        int i = ThreadID.get();
        flag[i] = false; // I'm not ready
    }
};
```





PetersonLock satisfies mutual exclusion

Suppose not. Consider the last execution of the lock method() by threads A and B . Inspecting the code, we see that

$$\begin{aligned} \text{write}_A(\text{flag}[A]=\text{true}) \rightarrow \text{write}_A(\text{victim}=A) \rightarrow \\ \text{read}_A(\text{flag}[B]) \rightarrow \text{read}_A(\text{victim}) \rightarrow CS_A \end{aligned} \quad (1)$$

$$\begin{aligned} \text{write}_B(\text{flag}[B]=\text{true}) \rightarrow \text{write}_B(\text{victim}=B) \rightarrow \\ \text{read}_B(\text{flag}[A]) \rightarrow \text{read}_B(\text{victim}) \rightarrow CS_B \end{aligned} \quad (2)$$

WLOG, assume that A was the last thread to write to the victim field

$$\text{write}_B(\text{victim}=B) \rightarrow \text{write}_A(\text{victim}=A) \quad (3)$$

Eq.(3) implies that A observed victim to be A in Eq. (1). Since A nevertheless entered its critical section, it must have observed $\text{flag}[B]$ to be *false*, so we have


$$\text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B]=\text{false}) \quad (4)$$

Eq.(2) ~ Eq.(4), together with the transitivity of precedence order, imply

$$\begin{aligned} \text{write}_B(\text{flag}[B]=\text{true}) \rightarrow \text{write}_B(\text{victim}=B) \rightarrow \\ \text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B]=\text{false}) \end{aligned}$$

It follows that $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow \text{read}_A(\text{flag}[B]=\text{false})$.

This is a contradiction because no other write to $\text{flag}[B]$ was performed before the critical section executions.





PetersonLock is starvation-free

Suppose not. WLOG, suppose that A runs forever in the `lock()` method. It must be executing the **while** statement, waiting until either `flag[B]` becomes *false* or `victim` is set to B .

What is B doing while A fails to make progress?

Perhaps B is repeatedly entering and leaving its critical section. If so, however, then B sets `victim` to B as soon as it reenters the critical section. Once `victim` is set to B , it does not change, and A must eventually return from the `lock()` method(), a contradiction.

So it must be that B is also stuck in the `lock()` method call, waiting until either `flag[A]` becomes *false* or `victim` is set to A . But `victim` cannot be both A and B , a contradiction. ☐





x86 Guaranteed Atomic Operations

Intel® 64 and IA-32 Architectures Software Developer's Manual Volume 3A: System Programming Guide, Part 1

8.1.1 Guaranteed Atomic Operations

The Intel486 processor (and newer processors since) guarantees that the following basic memory operations will always be carried out atomically:

- Reading or writing a byte
- Reading or writing a word aligned on a 16-bit boundary
- Reading or writing a doubleword aligned on a 32-bit boundary


The Pentium processor (and newer processors since) guarantees that the following additional memory operations will always be carried out atomically:

- Reading or writing a quadword aligned on a 64-bit boundary
- 16-bit accesses to uncached memory locations that fit within a 32-bit data bus

The P6 family processors (and newer processors since) guarantee that the following additional memory operation will always be carried out atomically:

- Unaligned 16-, 32-, and 64-bit accesses to cached memory that fit within a cache line

Accesses to cacheable memory that are split across cache lines and page boundaries are not guaranteed to be atomic by the Intel Core 2 Duo, Intel® Atom™, Intel Core Duo, Pentium M, Pentium 4, Intel Xeon, P6 family, Pentium, and Intel486 processors. The Intel Core 2 Duo, Intel Atom, Intel Core Duo, Pentium M, Pentium 4, Intel Xeon, and P6 family processors provide bus control signals that permit external memory subsystems to make split accesses atomic; however, nonaligned data accesses will seriously impact the performance of the processor and should be avoided.





x86 memory ordering guarantees

Intel® 64 and IA-32 Architectures Software Developer's Manual
Volume 3A: System Programming Guide, Part 1

CHAPTER 8 MULTIPLE-PROCESSOR MANAGEMENT

8.1	LOCKED ATOMIC OPERATIONS
8.1.1	Guaranteed Atomic Operations
8.1.2	Bus Locking
8.1.2.1	Automatic Locking
8.1.2.2	Software Controlled Bus Locking
8.1.3	Handling Self- and Cross-Modifying Code
8.1.4	Effects of a LOCK Operation on Internal Processor Caches
8.2	MEMORY ORDERING
8.2.1	Memory Ordering in the Intel® Pentium® and Intel486™ Processors
8.2.2	Memory Ordering in P6 and More Recent Processor Families
8.2.3	Examples Illustrating the Memory-Ordering Principles
8.2.3.1	Assumptions, Terminology, and Notation
8.2.3.2	Neither Loads Nor Stores Are Reordered with Like Operations
8.2.3.3	Stores Are Not Reordered With Earlier Loads
8.2.3.4	Loads May Be Reordered with Earlier Stores to Different Locations
8.2.3.5	Intra-Processor Forwarding Is Allowed
8.2.3.6	Stores Are Transitively Visible
8.2.3.7	Stores Are Seen in a Consistent Order by Other Processors
8.2.3.8	Locked Instructions Have a Total Order
8.2.3.9	Loads and Stores Are Not Reordered with Locked Instructions





Peterson Lock in action

[multithread_add_peterson.html](#)

```
class Peterson
{
    volatile bool X,Y;
    volatile pthread_t victim;

public:
    Peterson()
    {
        X = Y = false;
        victim = -1;
    }

    void lock()
    {
        if ( pthread_self()==tidX) { // thread X wanna to acquire lock
            X = true;
            victim = tidX;
            asm volatile("mfence" ::: "memory");
            while(Y && victim==tidX) ; // spin
        }
        else { //thread Y wanna acquire lock
            Y = true;
            victim = tidY;
            asm volatile("mfence" ::: "memory");
            while(X && victim==tidY) ; // spin
        }
    }
}
```

Prevent instruction reordering by processor

Prevent instruction reordering by optimizing compiler





Synchronization Hardware

Many systems provide hardware support for critical section code

Uniprocessors – could disable interrupts

Currently running code would execute without preemption

Generally too inefficient on multiprocessor systems

- ▶ Operating systems using this not broadly scalable

```
void lock()
{
    asm("cli"); // disable interrupt
}
```

```
void unlock()
{
    asm("sti"); // enable interrupt
}
```

Modern machines provide special atomic hardware instructions

- ▶ **Atomic = non-interruptable**

Either test memory word and set value

Or swap contents of two memory words

<http://en.wikipedia.org/wiki/Read-modify-write>





Solution to Critical-section Problem Using Locks

```
do {  
    acquire lock  
        critical section  
    release lock  
    remainder section  
} while (TRUE);
```





test_and_set Instruction

Definition:

```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

1. Executed atomically
2. Returns the original value of passed parameter
3. Set the new value of passed parameter to “TRUE”.





Solution using test_and_set()

Shared Boolean variable lock, initialized to FALSE

Solution:

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
        /* critical section */  
    lock = false;  
        /* remainder section */  
} while (true);
```





compare_and_swap Instruction

Definition:

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
  
    if (*value == expected)  
        *value = new_value;  
    return temp;  
}
```

1. Executed atomically
2. Returns the original value of passed parameter “value”
3. Set the variable “value” the value of the passed parameter “new_value” but only if “value” == “expected”. That is, the swap takes place only under this condition.





Solution using compare_and_swap

Shared integer “lock” initialized to 0;

Solution:

```
do {  
    while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
    /* critical section */  
    lock = 0;  
    /* remainder section */  
} while (true);
```






xchg (swap) on x86

```
*C:\Users\Hank\Downloads\postgresql-9.2.1.tar\postgresql-9.2.1\postgresql-9.2.1\src\include\storage\s_lock.h - Notepad++
File Edit Search View Encoding Language Settings Macro Run Plugins Window ?
s_lock.h
194
195 #ifdef __x86_64__ /* AMD Opteron, Intel EM64T */
196 #define HAS_TEST_AND_SET
197
198 typedef unsigned char slock_t;
199
200 #define TAS(lock) tas(lock)
201
202 static __inline__ int
203 tas(volatile slock_t *lock)
204 {
205     register slock_t _res = 1;
206
207     /*
208      * On Opteron, using a non-locking test before the locking instruction
209      * is a huge loss. On EM64T, it appears to be a wash or small loss,
210      * so we needn't bother to try to distinguish the sub-architectures.
211      */
212     __asm__ __volatile__ (
213         "    lock          \n"
214         "    xchgb    %0,%1 \n"
215         : "+q"(_res), "+m"(*lock)
216         :
217         : "memory", "cc");
218     return (int) _res;
219 }
220
221 #define SPIN_DELAY() spin_delay()
222
```

C++ source file length: 26285 lines: 1027 Ln: 214 Col: 18 Sel: 0 UNIX ANSI INS





Bounded-waiting Mutual Exclusion with TestandSet()

```
do {  
    waiting[i] = TRUE;  
    key = TRUE;  
    while (waiting[i] && key)  
        key = TestAndSet(&lock);  
    waiting[i] = FALSE; ← Wait until released or no one busy.  
    // critical section  
    j = (i + 1) % n;  
    while ((j != i) && !waiting[j])  
        j = (j + 1) % n; ← Look for a waiting process.  
    if (j == i)  
        lock = FALSE; ← No process waiting.  
    else  
        waiting[j] = FALSE; ← Process j is waiting. Release it.  
    // remainder section  
} while (TRUE);
```





Mutex Locks

Previous solutions are complicated and generally inaccessible to application programmers

OS designers build software tools to solve critical section problem

Simplest is mutex lock

Protect a critical section by first **acquire()** a lock then **release()** the lock

Boolean variable indicating if lock is available or not

Calls to **acquire()** and **release()** must be atomic

Usually implemented via hardware atomic instructions

But this solution requires **busy waiting**

This lock therefore called a **spinlock**





acquire() and release()

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;;  
}  
release() {  
    available = true;  
}  
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```





Semaphore

Synchronization tool that provides more sophisticated ways (than Mutex locks) for process to synchronize their activities.

Semaphore **S** – integer variable

Invented by Edsger Dijkstra in 1965

Can only be accessed via two indivisible (atomic) operations

wait() and **signal()**

- ▶ Originally called **P()** and **V()**

Definition of the **wait()** operation

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

Definition of the **signal()** operation

```
signal(S) {  
    S++;  
}
```





Semaphore Usage

Counting semaphore – integer value can range over an unrestricted domain

Binary semaphore – integer value can range only between 0 and 1

Same as a **mutex lock**

Can solve various synchronization problems

Consider P_1 and P_2 that require S_1 to happen before S_2

Create a semaphore “**synch**” initialized to 0

P1 :

$S_1;$


signal (synch) ;

P2 :

wait (synch) ;

$S_2;$

Can implement a counting semaphore **S** as a binary semaphore





Semaphore Implementation

Must guarantee that no two processes can execute the **wait()** and **signal()** on the same semaphore at the same time

Thus, the implementation becomes the critical section problem where the **wait** and **signal** code are placed in the critical section

Could now have **busy waiting** in critical section implementation

- ▶ But implementation code is short
- ▶ Little busy waiting if critical section rarely occupied

Note that applications may spend lots of time in critical sections and therefore this is not a good solution





Semaphore Implementation

Implementation of wait and signal

```
wait (S) {  
    while S <= 0  
        yield(); // give up CPU time slices  
    S--;  
}  
signal (S) {  
    S++;  
}
```

Both methods are atomic

The execution of wait(S) cannot be interrupted by other executions of wait(S) or executions of signal(S)

The execution of signal(S) cannot be interrupted by other executions of wait(S) or executions of signal(S)

Interrupts can occur during executions of wait(S) and signal(S)
(preemptive scheduling still works !)





Semaphore Implementation

With each semaphore there is an associated waiting queue.
Each entry in a waiting queue has two data items:

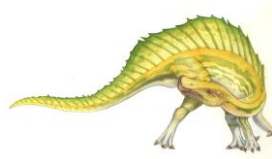
- value (of type integer)

- pointer to next record in the list

Two operations:

- block** – place the process invoking the operation on the appropriate waiting queue.

- wakeup** – remove one of processes in the waiting queue and place it in the scheduler ready queue.





Semaphore Implementation

Implementation of wait:

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}
```


The implementation is correct only if the two methods are atomic !

Implementation of signal:

```
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

How to ensure atomicity?






```
typedef struct __lock_t {
    int flag;
    int guard;
    queue_t *q;
} lock_t;


void lock_init(lock_t *m) {
    m->flag = 0;
    m->guard = 0;
    queue_init(m->q);
}

void lock(lock_t *m) {
    while (TestAndSet(&m->guard, 1) == 1)    ; //acquire guard lock by spinning
    if (m->flag == 0) {
        m->flag = 1; // lock is acquired
        m->guard = 0;
    } else {
        queue_add(m->q, getpid());
        m->guard = 0;
        block();
    }
}

void unlock(lock_t *m) {
    while (TestAndSet(&m->guard, 1) == 1)    ; //acquire guard lock by spinning
    if (queue_empty(m->q))
        m->flag = 0; // let go of lock; no one wants it
    else
        wakeup(queue_remove(m->q)); // hold lock (for next thread!)
    m->guard = 0;
}
```

wakeup/waiting race





Deadlock and Starvation

Deadlock – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

Let S and Q be two semaphores initialized to 1

P_0
`wait(S) ;`
`wait(Q) ;`
`...`
`signal(S) ;`
`signal(Q) ;`

P_1
`wait(Q) ;`
`wait(S) ;`
`...`
`signal(Q) ;`
`signal(S) ;`

Starvation – **indefinite blocking**

A process may never be removed from the semaphore queue in which it is suspended

Priority Inversion – Scheduling problem when lower-priority process holds a lock needed by higher-priority process

Solved via **priority-inheritance protocol**





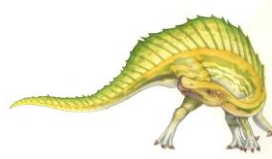
Classical Problems of Synchronization

Classical problems used to test newly-proposed synchronization schemes

- Bounded-Buffer Problem

- Readers and Writers Problem

- Dining-Philosophers Problem





Bounded-Buffer Problem

N buffers, each can hold one item

Semaphore **full** initialized to the value 0

Semaphore **empty** initialized to the value N .






Bounded Buffer Problem (Cont.)

The structure of the producer process

```
do {  
  
    // produce an item  
    wait (empty);  
    // add the item to the buffer  
    signal (full);  
} while (TRUE);
```

The structure of the consumer process

```
do {  
    wait (full);  
    // remove an item from buffer  
    signal (empty);  
    // consume the item  
} while (TRUE);
```





Bounded-Buffer Problem

N buffers, each can hold one item

Semaphore **mutex** initialized to the value 1

Semaphore **full** initialized to the value 0

Semaphore **empty** initialized to the value N .






Bounded Buffer Problem (Cont.)

The structure of the producer process

```
do {  
  
    // produce an item  
  
    wait (empty);  
    wait (mutex);  
  
    // add the item to the buffer  
  
    signal (mutex);  
    signal (full);  
} while (TRUE);
```

The structure of the consumer process

```
do {  
    wait (full);  
    wait (mutex);  
  
    // remove an item from buffer  
  
    signal (mutex);  
    signal (empty);  
    // consume the item  
} while (TRUE);
```





Readers-Writers Problem

A data set is shared among a number of concurrent processes

Readers – only read the data set; they do **not** perform any updates

Writers – can both read and write

Problem – allow multiple readers to read at the same time

Only one single writer can access the shared data at the same time

Several variations of how readers and writers are considered – all involve some form of priorities

Shared Data

Data set

Semaphore **rw_mutex** initialized to 1

Semaphore **mutex** initialized to 1

Integer **read_count** initialized to 0





Readers-Writers Problem (Cont.)

The structure of a writer process

```
do {  
    wait(rw_mutex);  
    ...  
    /* writing is performed */  
    ...  
    signal(rw_mutex);  
} while (true);
```





Readers-Writers Problem (Cont.)

The structure of a reader process

```
do {  
    wait(mutex);  
    read_count++;  
    if (read_count == 1)  
        wait(rw_mutex);  
    signal(mutex);  
  
    ...  
    /* reading is performed */  
    ...  
    wait(mutex);  
    read_count--;  
    if (read_count == 0)  
        signal(rw_mutex);  
    signal(mutex);  
} while (true);
```





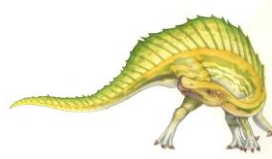
Readers-Writers Problem Variations

First variation – no reader kept waiting unless writer has permission to use shared object

Second variation – once writer is ready, it performs the write ASAP


Both may have starvation leading to even more variations

Problem is solved on some systems by kernel providing reader-writer locks





Dining-Philosophers Problem



Philosophers spend their lives alternating thinking and eating

Don't interact with their neighbors, occasionally try to pick up 2 chopsticks (one at a time) to eat from bowl

Need both to eat, then release both when done

In the case of 5 philosophers

Shared data

- ▶ Bowl of rice (data set)
- ▶ Semaphore **chopstick** [5] initialized to 1





Dining-Philosophers Problem Algorithm

The structure of Philosopher *i*:

```
do {  
    wait (chopstick[i] );  
    wait (chopStick[ (i + 1) % 5] );  
  
    // eat  
  
    signal (chopstick[i] );  
    signal (chopstick[ (i + 1) % 5] );  
  
    // think  
  
} while (TRUE);
```

What is the problem with this algorithm?





Dining-Philosophers Problem Algorithm (Cont.)

Deadlock handling

Allow at most 4 philosophers to be sitting simultaneously at the table.

Allow a philosopher to pick up the forks only if both are available (picking must be done in a critical section).

Use an asymmetric solution -- an odd-numbered philosopher picks up first the left chopstick and then the right chopstick. Even-numbered philosopher picks up first the right chopstick and then the left chopstick.





Problems with Semaphores

Incorrect use of semaphore operations:

signal (mutex) wait (mutex)

wait (mutex) ... wait (mutex)

Omitting of wait (mutex) or signal (mutex) (or both)

Deadlock and starvation are possible.





Problems with Semaphores

Used for 2 independent purposes

- Mutual exclusion

- Condition synchronization

Hard to get right

- signal (mutex) wait (mutex)

- wait (mutex) ... wait (mutex)

- Omitting of wait (mutex) or signal (mutex) (or both)

- Small mistake easily leads to deadlock / livelock

Would it be nice to have?

- Separation of mutual exclusion and condition synchronization

- Automatic *wait* and *signal*





Monitors


Invented by Tony Hoare in 1974

Like a C++ class

Consists of *vars* and *procedures*

3 key differences from a regular class:

- ▶ Only one thread in a monitor at a time (automatic mutual exclusion)
- ▶ Special type of variable, called “condition variable”
 - 3 special ops on a condition variable: *wait*, *signal*, and *broadcast*
- ▶ No public variables allowed (must call procedures to access variables)





Monitors

A high-level abstraction that provides a convenient and effective mechanism for process synchronization

Only one process may be active within the monitor at a time

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { .... }
    ...


    procedure Pn (...) {.....}

    Initialization code ( ....) { ... }
    ...
}
```





Schematic view of a Monitor





Bounded Buffer by Monitor

```
BoundedBuffer {  
  
    int BUFFER[MAX_SIZE];  
    int head, tail, size;  
  
    Enque (int v) {  
        BUFFER[tail] = v;  
        tail = (tail+1) % MAX_SIZE;  
        size++;  
    }  
  
    Deque (int v) {  
        int i = head;  
        head = (head+1) % MAX_SIZE;  
        size--;  
        return BUFFER[i];  
    }  
  
    Init () {  
        head = tail = size = 0;  
    }  
};
```

Any problem?





Bounded Buffer by Monitor

```
BoundedBuffer {  
  
    int BUFFER[MAX_SIZE];  
    int head, tail, size;  
  
    Enque (int v) {  
        while( size == MAX_SIZE);  
        BUFFER[tail] = v;  
        tail = (tail+1) % MAX_SIZE;  
        size++;  
    }  
  
    Deque (int v) {  
        while(size==0);  
        int i = head;  
        head = (head+1) % MAX_SIZE;  
        size--;  
        return BUFFER[i];  
    }  
}
```

Any problem?





Condition Variables

Need a mechanism for condition synchronization

condition x, y;

Automatic *unlock* and *lock* for mutual exclusion

Two operations on a condition variable:

`cond.wait ()`

- ▶ Thread is put on queue for “cond”, goes to sleep.

`cond.signal ()`

- ▶ If queue for “cond” not empty, wake up on thread

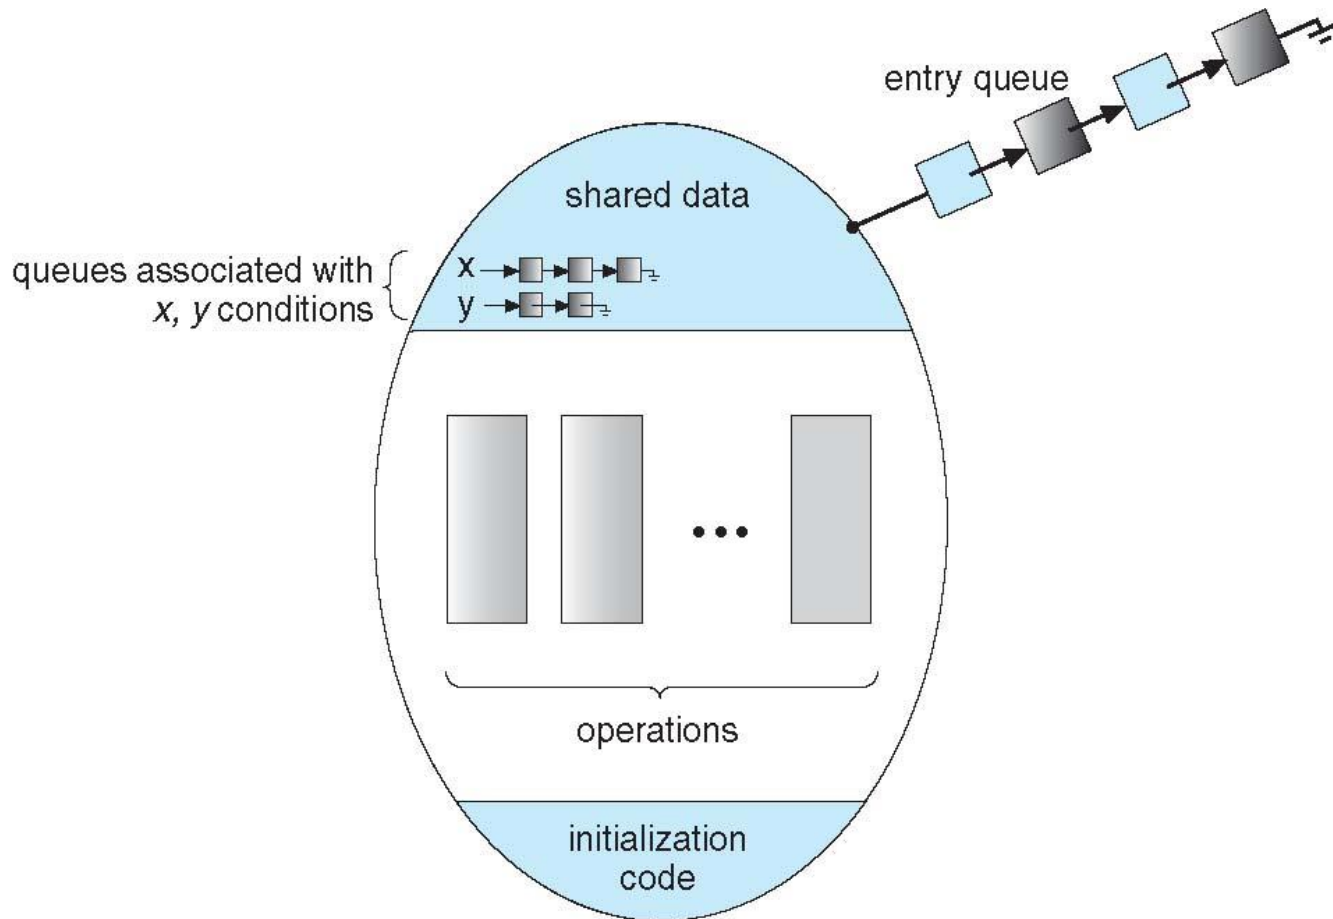
`cond.broadcast()`

- ▶ Wake up all threads waiting on queue for “cond”





Monitor with Condition Variables





Semantics of Signal

Signal and Wait (Hoare-style)

Signaler passes lock, CPU to waiter; waiter runs immediately

Waiter gives lock, CPU back to signaler when

- ▶ It exits critical section
- ▶ Or, it waits again

Signal and Continue (Mesa-style)

signaler continues executing

waiter put on ready queue

when waiter actually gets to run

- ▶ May have to wait for lock again
- ▶ State may have changed! Use “while”, not “if”

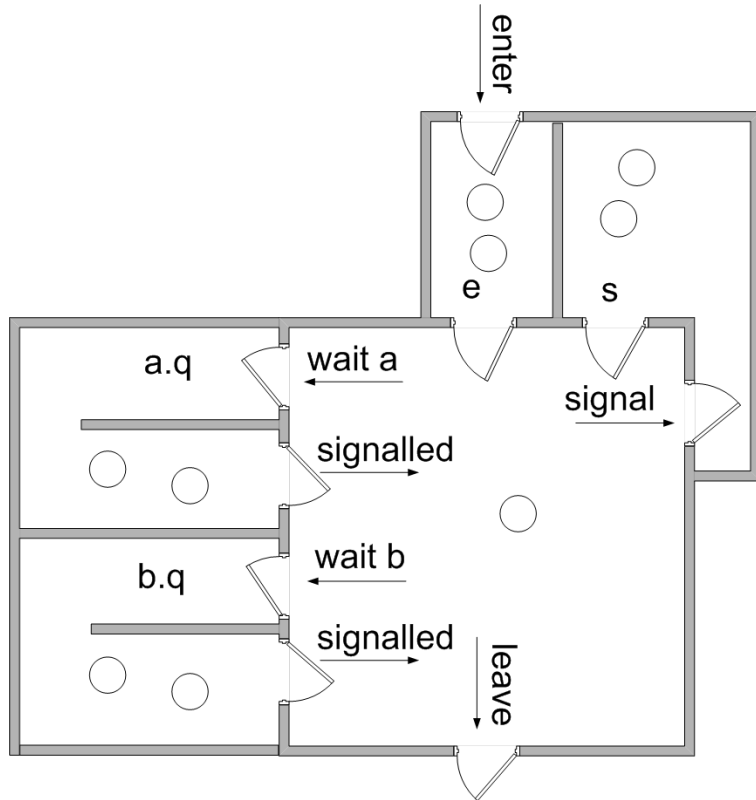
Used in Java, Pthread, ...)

<http://www.cs.mtu.edu/~shene/NSF-3/e-Book/MONITOR/monitor-types.html>






Hoare Style vs Mesa Style



Hoare Style



Mesa Style





Bounded Buffer by Monitor

```
BoundedBuffer {
```

```
    int BUFFER[MAX_SIZE];
```

```
    int head, tail, size;
```

```
    cond full, empty;
```

```
    Enqueue (int v) {
```

```
        while( size == MAX_SIZE)
```

```
            full.wait();
```

```
        BUFFER[tail] = v;
```

```
        tail = (tail+1) % MAX_SIZE;
```

```
        size++;
```

```
        if (size ==1) empty.signal();
```

```
    }
```

```
    Dequeue (int v) {
```

```
        while(size==0)
```

```
            empty.wait();
```

```
        int i = head;
```

```
        head = (head+1) % MAX_SIZE;
```

```
        size--;
```

```
        if ( size == MAX_SIZE-1)
```

```
            full.signal();
```

```
        return BUFFER[i];
```

```
    }
```

```
    .....
```

```
}
```





Monitor Solution to Dining Philosophers

```
monitor DiningPhilosophers
{
    enum { THINKING; HUNGRY, EATING} state [5] ;
    condition self [5];

    void pickup (int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING) self[i].wait;
    }

    void putdown (int i) {
        state[i] = THINKING;
        // test left and right neighbors
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
}
```





Solution to Dining Philosophers (Cont.)

```
void test (int i) {
    if ((state[(i + 4) % 5] != EATING) &&
        (state[i] == HUNGRY) &&
        (state[(i + 1) % 5] != EATING) ) {
        state[i] = EATING ;
        self[i].signal () ;
    }
}

initialization_code() {
    for (int i = 0; i < 5; i++)
        state[i] = THINKING;
}
}
```





Solution to Dining Philosophers (Cont.)

Each philosopher i invokes the operations **pickup()** and **putdown()** in the following sequence:

DiningPhilosophers.pickup(i) ;

EAT

DiningPhilosophers.putdown(i) ;

No deadlock, but starvation is possible





Hoare Style Monitor Implementation (using semaphores)

Need mutual exclusion semaphore **mutex** (init to 1) so that only one process is active within monitor

Need a semaphore **next** (**next to exit**) for the signaling process to suspend itself

initialized to **zero**

next_count is number of processes blocked on **next**

Before exiting a procedure, process must either:

Signal other waiting processes in monitor **next** before exiting, or

Signal **mutex** and exit





Monitor Implementation (Hoare Style)

The monitor “compiler” has to automatically insert this code into compiled procedures:

Procedure F:

```
wait(mutex);
```

```
...
```

```
body of F
```

```
...
```

```
if (next_count>0)
```

```
    signal(next);
```

```
else
```

```
    signal(mutex);
```

```
end;
```





Condition Variable Implementation (Hoare)

Each condition x has a count, and a standard semaphore (with associated queue) **initialized to 0**

```
x.wait() {  
    x.count++;  
    if (next_count > 0)  
        signal(next);  
    else  
        signal(mutex);  
    wait(x.sem);  
    x.count--;  
}
```

```
x.signal() {  
    if (x.count > 0){  
        next_count++;  
        signal(x.sem);  
        wait(next);  
        next_count--;  
    }  
}
```





Monitor Implementation (Mesa Style)

Need mutual exclusion semaphore **mutex** (init to 1) so that only one process is active within monitor

The monitor “compiler” has to automatically insert this code into compiled procedures:

Procedure F:

wait(mutex);

...

body of F

...

signal(mutex);

end;





Condition Variable Implementation (Mesa)

Each condition x has a count, and a standard semaphore (with associated queue) **initialized to 0**

```
x.wait() {  
    x.count++;  
    signal(mutex);  
    wait(x.sem);  
    wait(mutex);  
}
```

```
x.signal() {  
    if (x.count > 0){  
        x.count--;  
        signal(x.sem);  
    }  
}
```





Difference between Monitors and Semaphores

Monitors enforce mutual exclusion

Semaphore wait vs. Monitor wait

Semaphore wait blocks if value is 0

Monitor wait always blocks

Semaphore signal vs Monitor signal

Semaphore signal either wakes up a thread or increments value

Monitor signal only has effect if a thread waiting

Semaphores have “memory”





Resuming Processes within a Monitor

If several processes queued on condition x , and $x.\text{signal}()$ executed, which should be resumed?

FCFS frequently not adequate

conditional-wait construct of the form $x.\text{wait}(c)$

Where c is **priority number**

Process with lowest number (highest priority) is scheduled next





Single Resource allocation

Allocate a single resource among competing processes using priority numbers that specify the maximum time a process plans to use the resource

```
R.acquire (t) ;  
    ...  
    access the resource ;  
    ...  
  
R.release ;
```

Where R is an instance of type **ResourceAllocator**





A Monitor to Allocate Single Resource

```
monitor ResourceAllocator
{
    boolean busy;
    condition x;

    void acquire(int time) {
        if (busy)
            x.wait(time);
        busy = TRUE;
    }

    void release() {
        busy = FALSE;
        x.signal();
    }

    initialization code() {
        busy = FALSE;
    }
}
```





Synchronization Examples

Solaris

Windows

Linux

Pthreads





Solaris Synchronization

Implements a variety of locks to support multitasking, multithreading (including real-time threads), and multiprocessing

Uses **adaptive mutexes** for efficiency when protecting data from short code segments

- Starts as a standard semaphore spin-lock

- If lock held, and by a thread running on another CPU, spins

- If lock held by non-run-state thread, block and sleep waiting for signal of lock being released

Uses **condition variables**

Uses **readers-writers** locks when longer sections of code need access to data

Uses **turnstiles** to order the list of threads waiting to acquire either an adaptive mutex or reader-writer lock

- Turnstiles are per-lock-holding-thread, not per-object

Priority-inheritance per-turnstile gives the running thread the highest of the priorities of the threads in its turnstile





Windows Synchronization

Uses interrupt masks to protect access to global resources on uniprocessor systems

Uses **spinlocks** on multiprocessor systems

Spinlocking-thread will never be preempted

Also provides **dispatcher objects** user-land which may act mutexes, semaphores, events, and timers

Events

- ▶ An event acts much like a condition variable

Timers notify one or more thread when time expired

Dispatcher objects either **signaled-state** (object available) or **non-signaled state** (thread will block)





Linux Synchronization

Linux:

Prior to kernel Version 2.6, disables interrupts to implement short critical sections

Version 2.6 and later, fully preemptive

Linux provides:

Semaphores

atomic integers

spinlocks

reader-writer versions of both

On single-cpu system, spinlocks replaced by enabling and disabling kernel preemption





Pthreads Synchronization

Pthreads API is OS-independent

It provides:

- mutex locks

- condition variable

Non-portable extensions include:

- read-write locks

- spinlocks





Java-style monitors

Integrated into the class mechanism

Annotation “synchronized” can be applied to a member function

- ▶ This function executes with implicit mutual exclusion

Wait, Signal, and Broadcast are called monitor wait, notify, and notifyAll, respectively

<http://docs.oracle.com/javase/tutorial/essential/concurrency/syncmeth.html>

```
public class SynchronizedCounter {  
    private int c = 0;  
  
    public synchronized void increment() {  
        c++;  
    }  
  
    public synchronized void decrement() {  
        c--;  
    }  
  
    public synchronized int value() {  
        return c;  
    }  
}
```





Alternative Approaches

Transactional Memory

OpenMP

Functional Programming Languages





Transactional Memory

A **memory transaction** is a sequence of read-write operations to memory that are performed atomically.

```
void update()  
{  
    /* read/write memory */  
}
```





OpenMP

OpenMP is a set of compiler directives and API that support parallel programming.

```
void update(int value)
{
    #pragma omp critical
    {
        count += value
    }
}
```

The code contained within the `#pragma omp critical` directive is treated as a critical section and performed atomically.





Functional Programming Languages

Functional programming languages offer a different paradigm than procedural languages in that they do not maintain state.

Variables are treated as immutable and cannot change state once they have been assigned a value.

There is increasing interest in functional languages such as Erlang and Scala for their approach in handling data races.



End of Chapter 5

